

# Probabilistic Graphical Models

## Building Blocks of Deep Learning

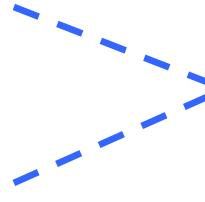
Zhiteng Hu

Lecture 16, March 18, 2019

Reading: see class homepage



# Overview: Deep Learning & Generative Models

- 3/6 Lecture 15 Statistical and Algorithmic Foundations of Deep Learning
- 3/18 Lecture 16 Building blocks of DL       *Model Architectures*
- 3/20 Lecture 17 Deep generative models (part 1)       *Inference & Learning*
- 3/25 Lecture 18 Deep generative models (part 2)
- 3/27 Lecture 19 A unified view of deep generative models  *Advanced Topics*





# Outline

- Convolutional Networks (ConvNets)
- Recurrent Networks (RNNs)
  - Long-range dependency, vanishing
  - LSTM
  - RNNs in different forms
- Attention Mechanisms
  - (Query, Key, Value)
  - Attention on Text and Images
- Transformers: Multi-head Attention
  - Transformer
  - BERT

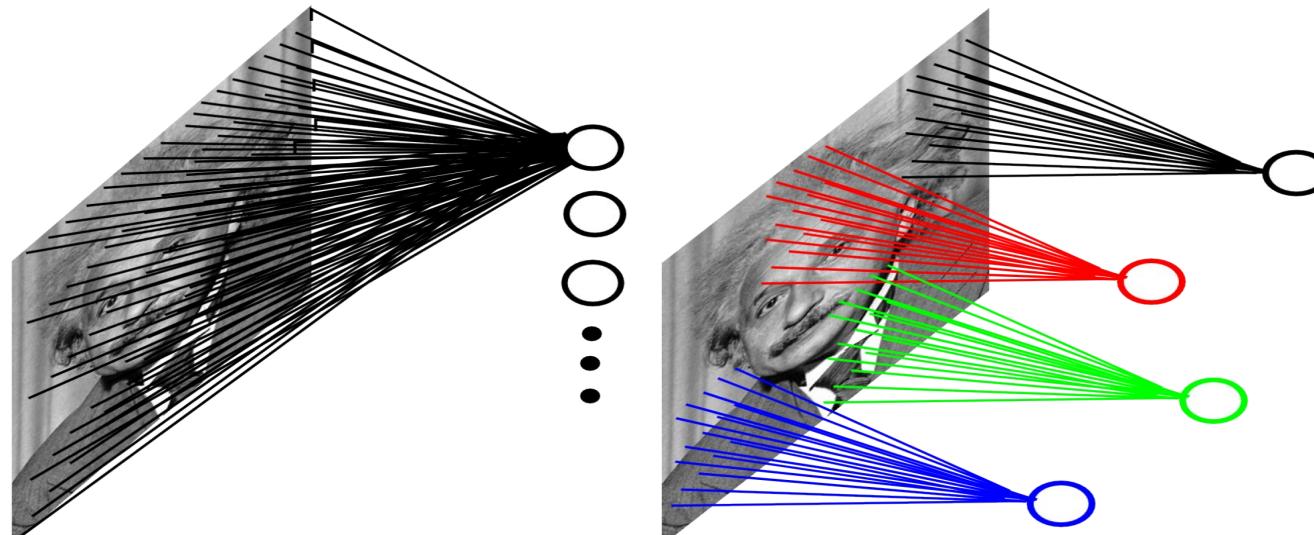




# Convolutional Networks (ConvNets)

- Biologically-inspired variants of MLPs [LeCun et al. NIPS 1989]
  - Receptive field [Hubel & Wiesel 1962; Fukushima 1982]
    - Visual cortex contains a complex arrangement of cells
    - These cells are sensitive to small **sub-regions** of the visual field
  - The sub-regions are **tiled** to cover the entire visual field

Exploit the strong spatially local correlation present in natural images





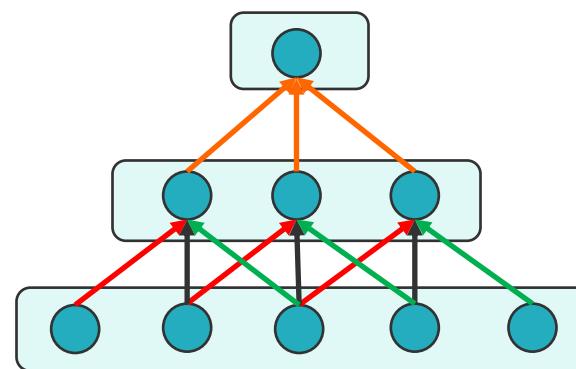
# Convolutional Networks (ConvNets)

- Sparse connectivity
- Shared weights
- Increasingly “global” receptive fields
  - simple cells detect local features
  - complex cells “pool” the outputs of simple cells within a retinotopic neighborhood.

**Feature maps  $m + 1$**

**Feature maps  $m$**

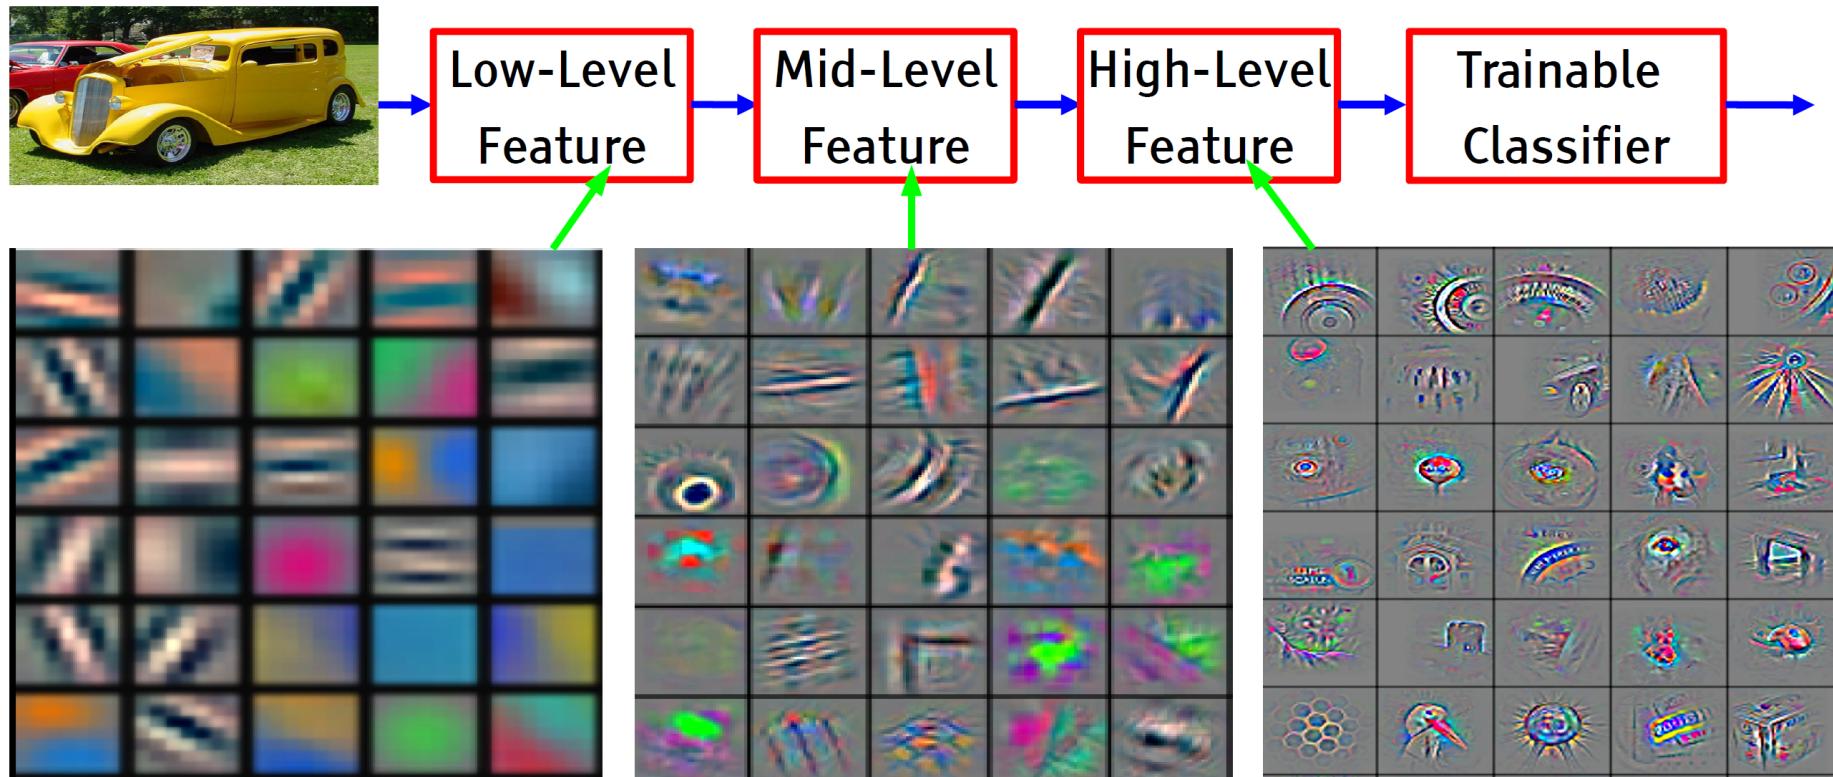
**Feature maps  $m - 1$**





# Convolutional Networks (ConvNets)

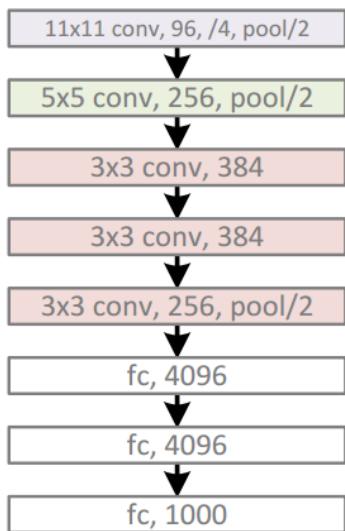
- Hierarchical Representation Learning [Zeiler & Fergus 2013]



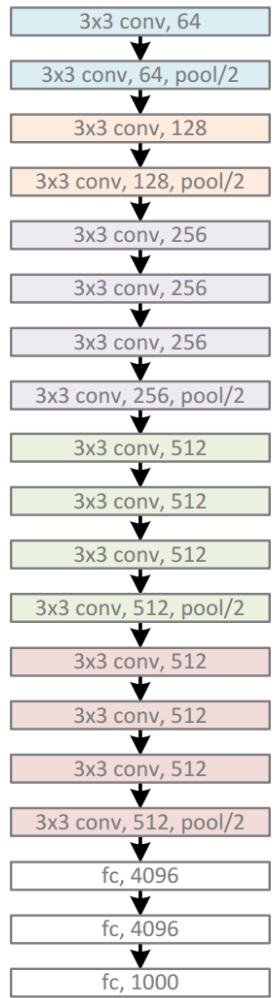


# Evolution of ConvNets

# AlexNet, 8 layers



# VGG, 19 layers



# GoogleNet, 22 layers



# ResNet, 152 layers





# Outline

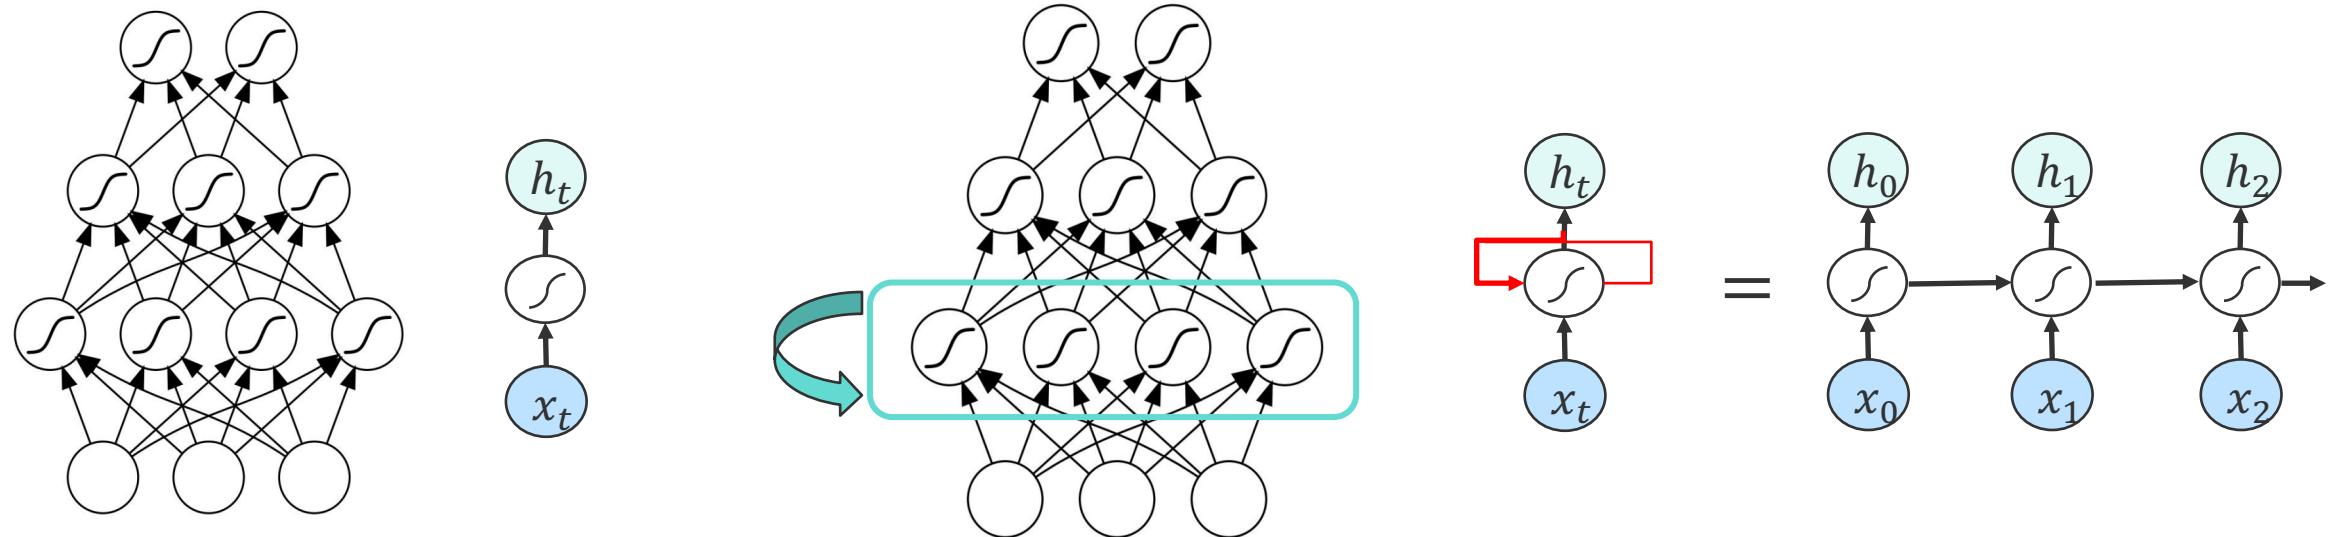
- Convolutional Networks (ConvNets)
- Recurrent Networks (RNNs)
  - Long-range dependency, vanishing
  - LSTM
  - RNNs in different forms
- Attention Mechanisms
  - (Query, Key, Value)
  - Attention on Text and Images
- Transformers: Multi-head Attention
  - Transformer
  - BERT





# ConvNets → Recurrent Networks (RNNs)

- Spatial Modeling *vs.* Sequential Modeling
- Fixed *vs.* variable number of computation steps.



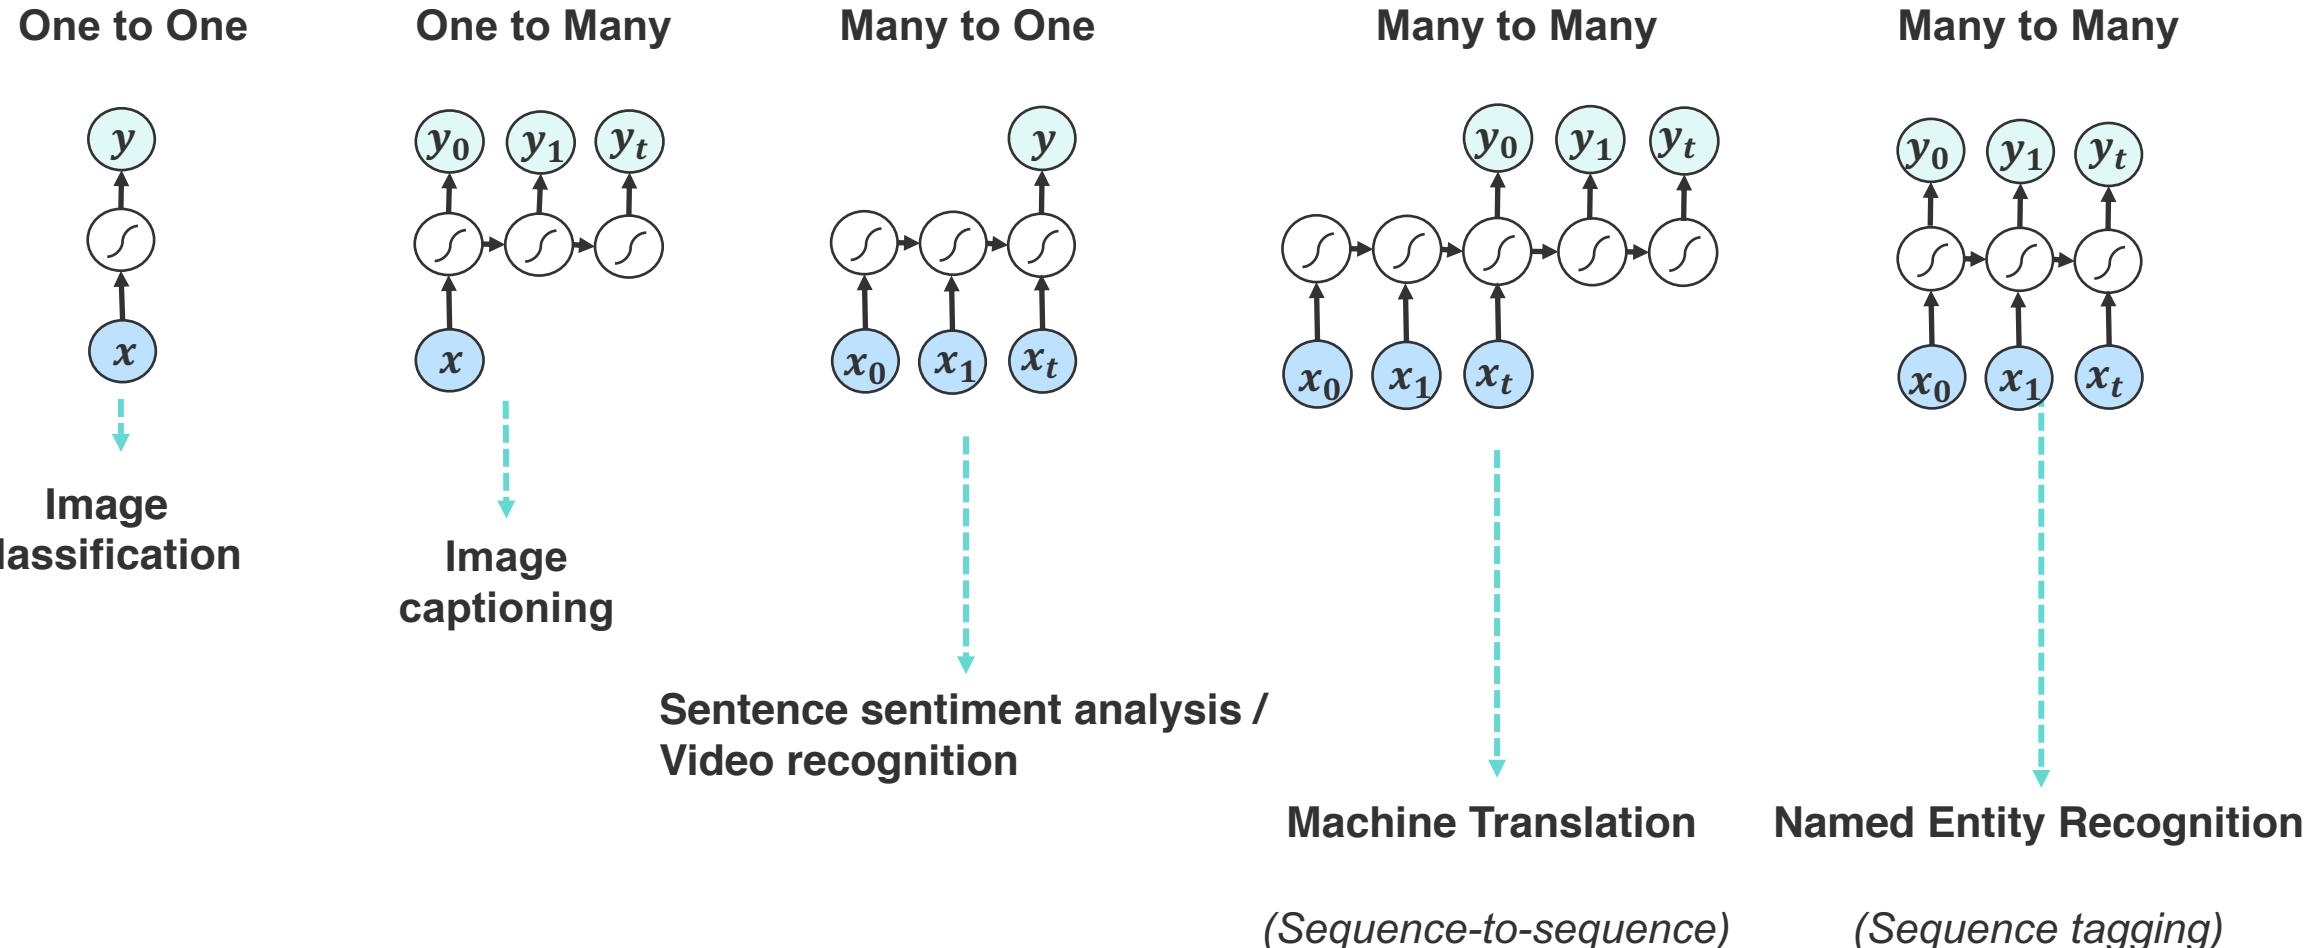
The output depends ONLY  
on the **current input**

The hidden layers and the output  
additionally depend on **previous states**  
of the hidden layers





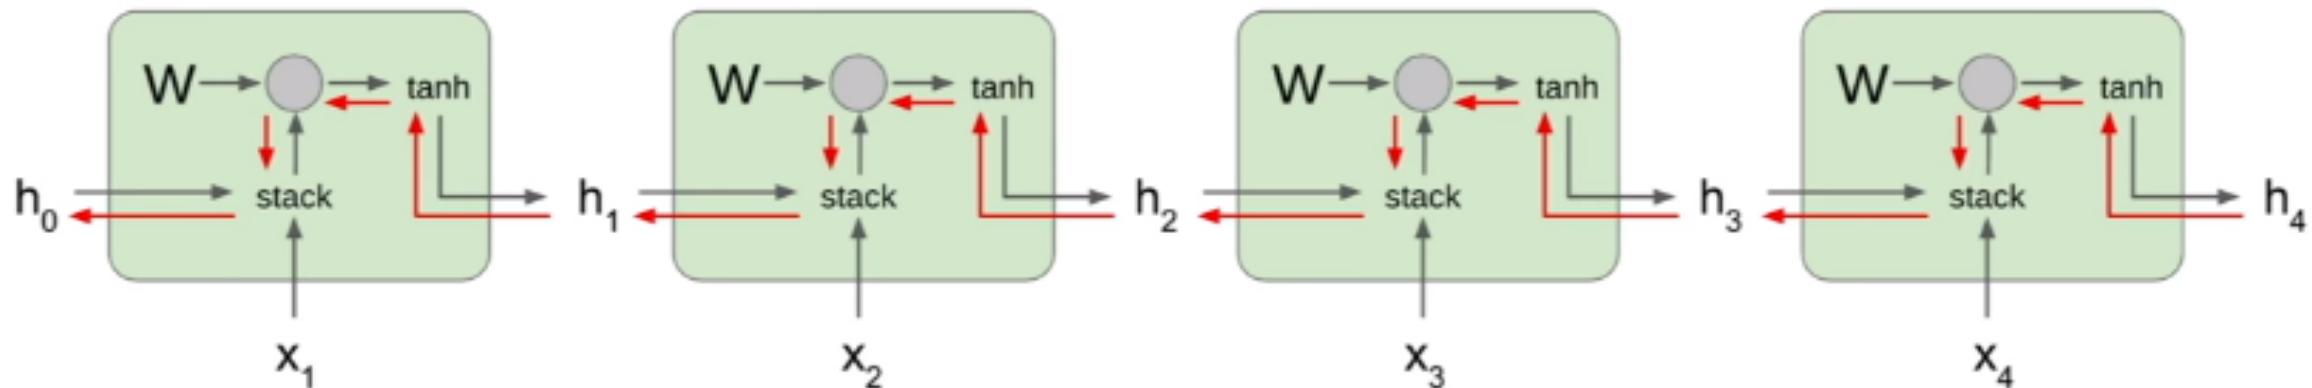
# RNNs in Various Forms





# Vanishing / Exploding Gradients in RNNs

$$h_t = \tanh(W^{hh}h_{t-1} + W^{hx}x_t)$$



Computing gradient of  $h_0$  involves many factors of  $W$  (and repeated  $\tanh$ )

Largest singular value  $> 1$ :  
**Exploding gradients**

Largest singular value  $< 1$ :  
**Vanishing gradients**

→ **Gradient clipping:** Scale gradient if its norm is too big

```
grad_norm = np.sum(grad * grad)
if grad_norm > threshold:
    grad *= (threshold / grad_norm)
```

Bengio et al., 1994 “Learning long-term dependencies with gradient descent is difficult”  
Pascanu et al., 2013 “On the difficulty of training recurrent neural networks”

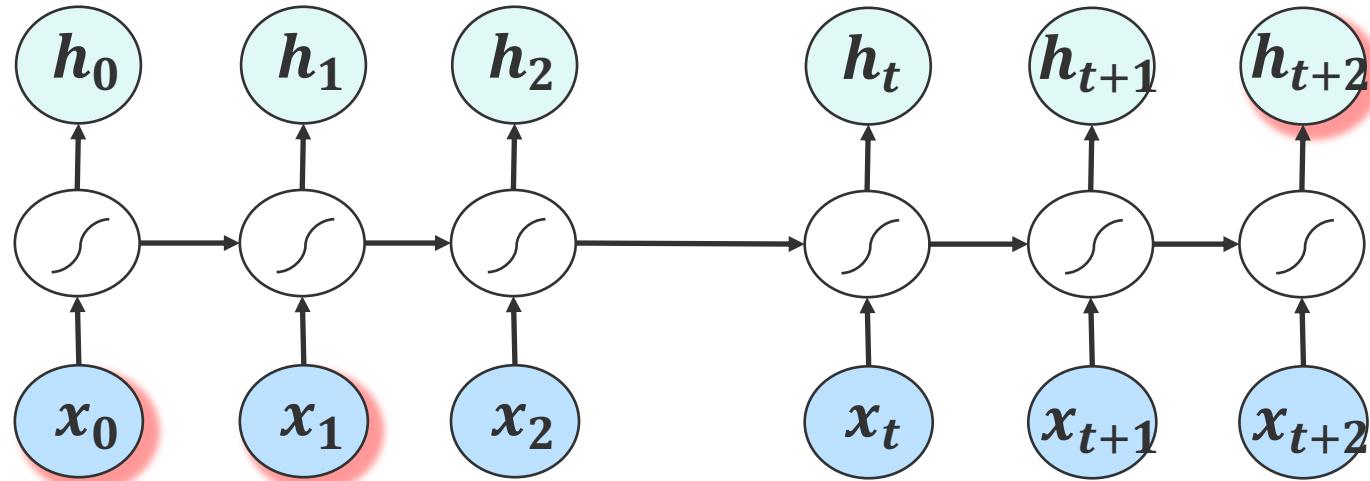
© Eric Xing @ CMU, 2005-2019

11





# Long-term Dependency Problem



I live in France and I know French

I live in France, a beautiful country, and I know French

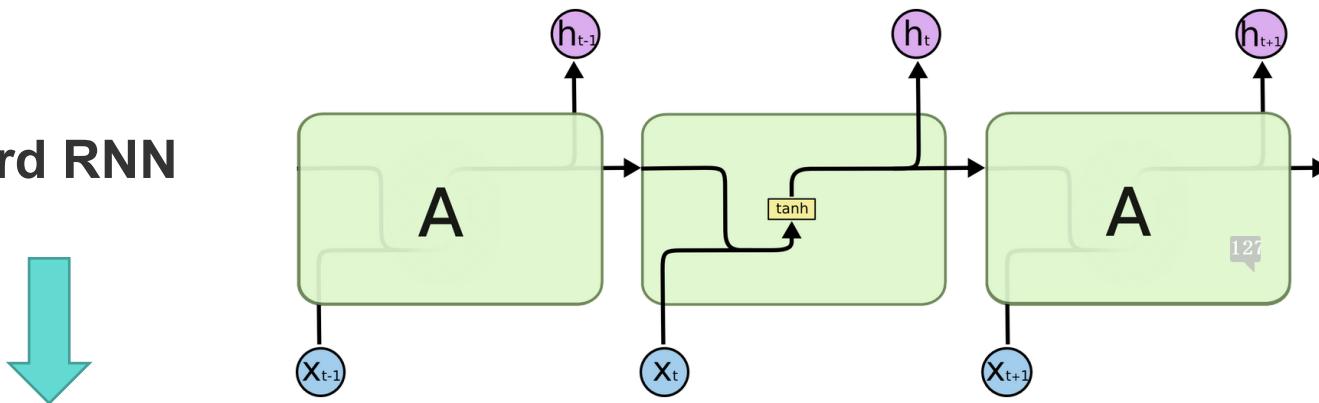




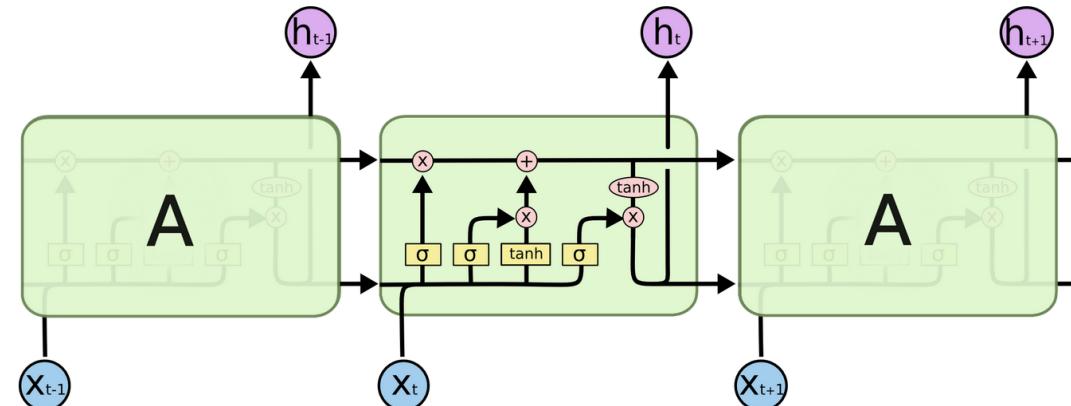
# Long Short Term Memory (LSTM)

- LSTMs are designed to explicitly alleviate the long-term dependency problem [Horchreiter & Schmidhuber (1997)]

Standard RNN



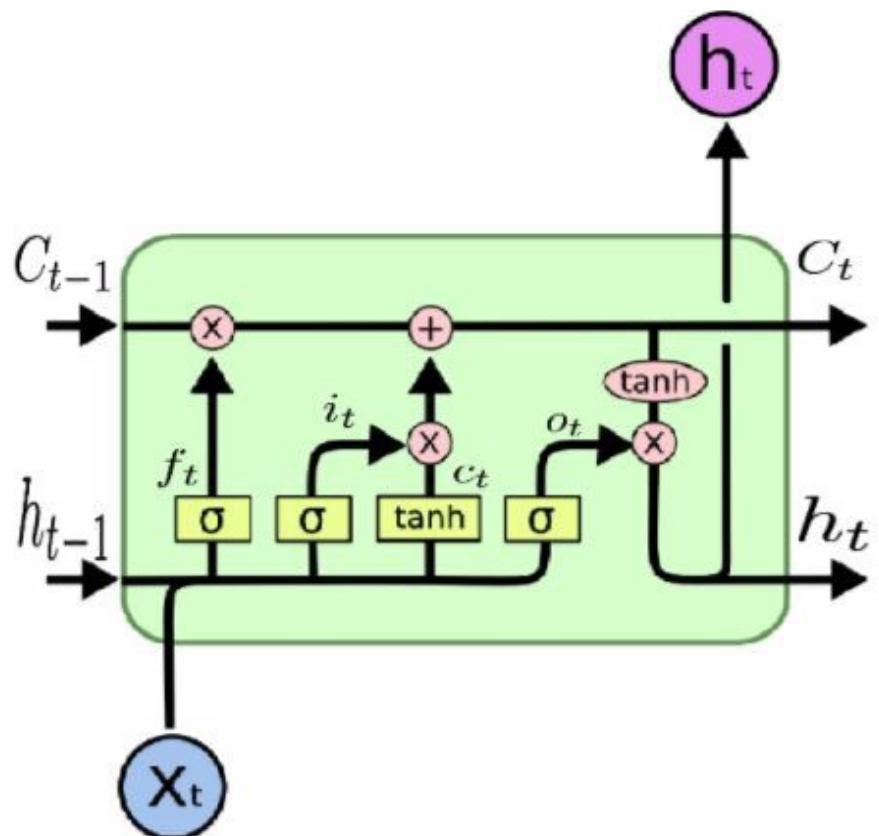
LSTM





# Long Short Term Memory (LSTM)

- Linear memory cells + multiplicative gate units to store read, write, and reset information



Gating functions to select information for passing and remembering

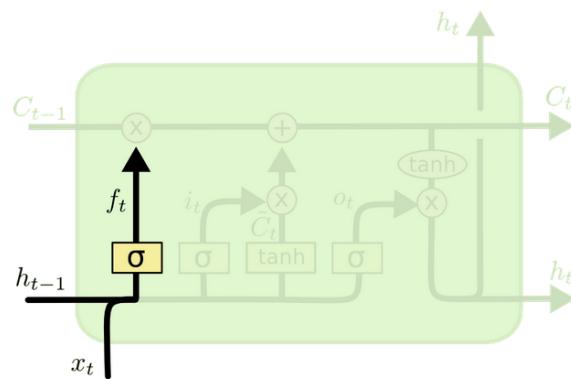
- Forget gate: whether to erase cell (reset)
- Input gate: whether to write to cell (write)
- Output gate: how much to reveal cell (read)





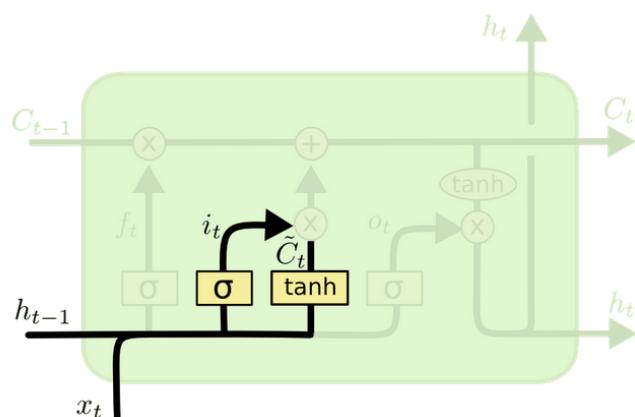
# Long Short Term Memory (LSTM)

- Forget gate: decides what must be removed from  $h_{t-1}$



$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

- Input gate: decides what new information to store in the cell



$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

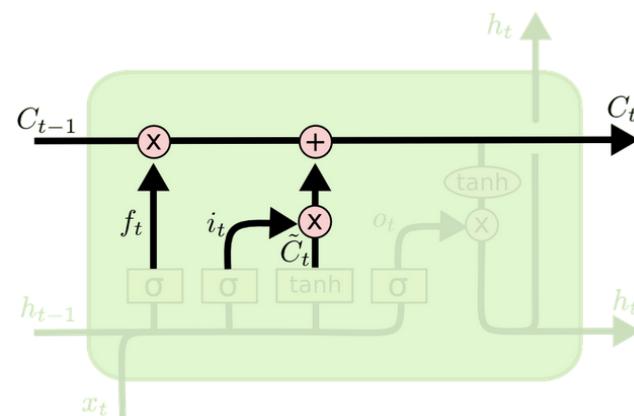
$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$





# Long Short Term Memory (LSTM)

- Update cell state:

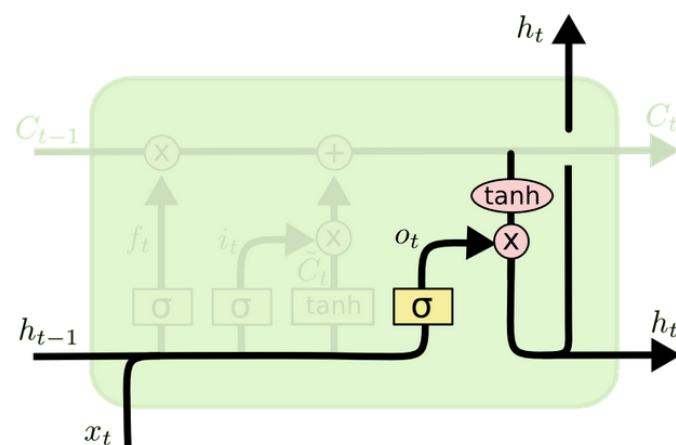


$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

forgetting unneeded things

scaling the new candidate values by how much we decided to update each state value.

- Output gate: decides what to output from our cell state



$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh(C_t)$$

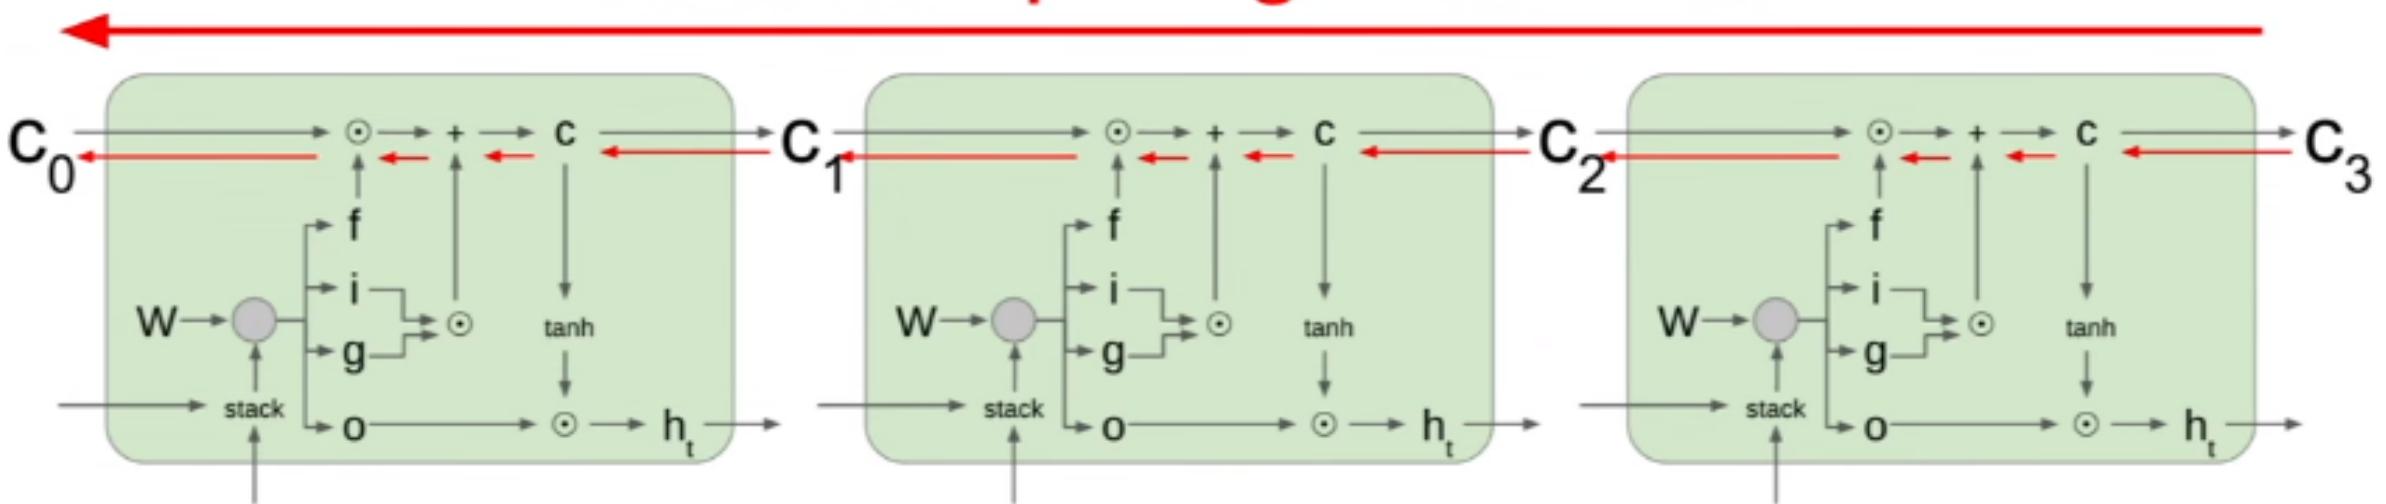
sigmoid decides what parts of the cell state we're going to output





# Backpropagation in LSTM

Uninterrupted gradient flow!



- No multiplication with matrix  $W$  during backprop
- Multiplied by different values of forget gate  $\rightarrow$  less prone to vanishing/exploding gradient





# RNNs in Various Forms

One to One



Image classification

One to Many

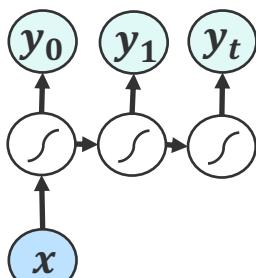
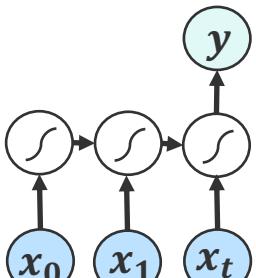


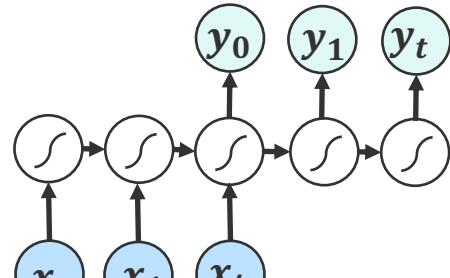
Image captioning

Many to One



Sentence sentiment analysis /  
Video recognition

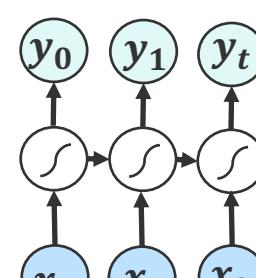
Many to Many



Machine Translation

(Sequence-to-sequence)

Many to Many



Named Entity Recognition

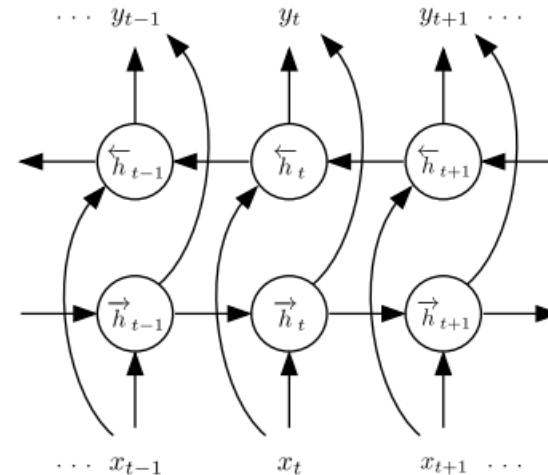
(Sequence tagging)



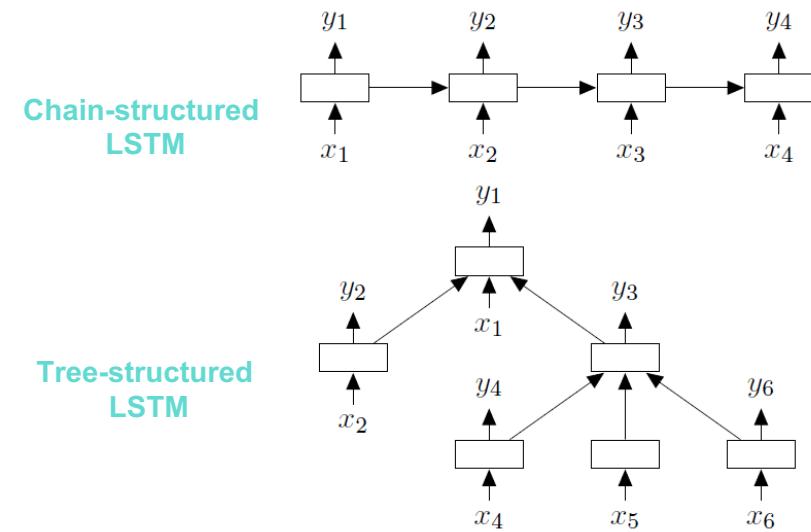


# RNNs in Various Forms

- Bi-directional RNN
  - Hidden state is the concatenation of both forward and backward hidden states.
  - Allows the hidden state to capture both **past** and **future** information.
- Tree-structured RNN
  - Hidden states condition on both an input vector and the hidden states of **arbitrarily** many child units.
  - Standard LSTM = a special case of tree-LSTM where each internal node has exactly one child.



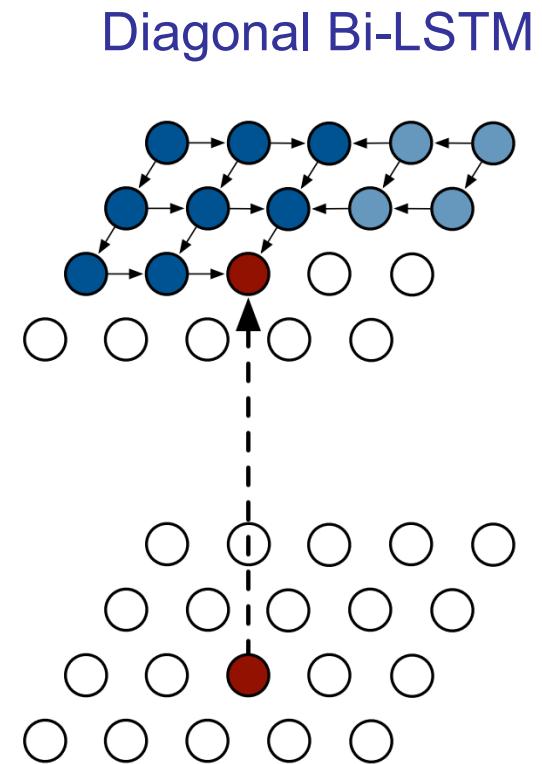
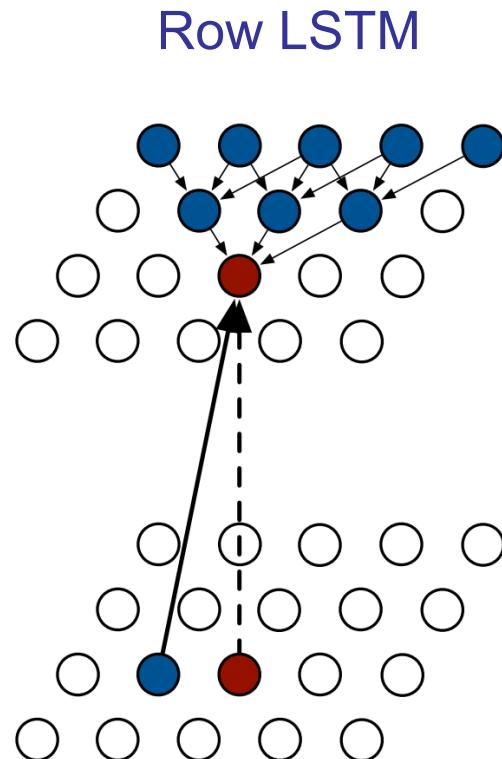
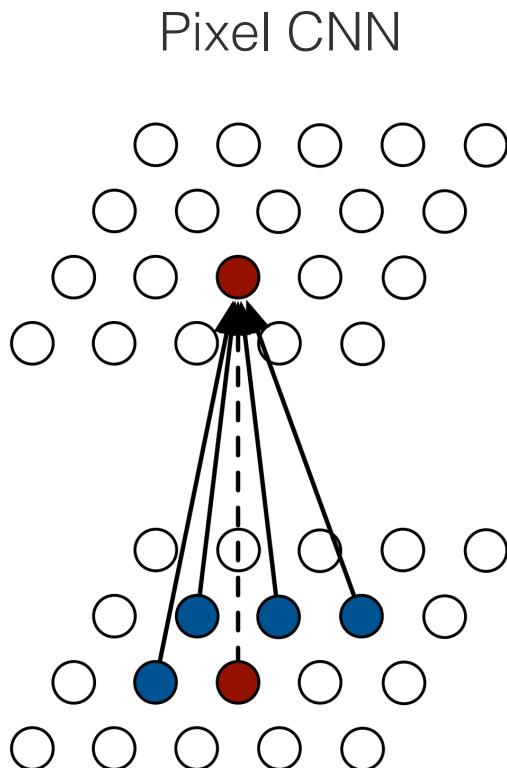
[Speech Recognition with Deep Recurrent Neural Networks, Alex Graves]





# RNNs in Various Forms

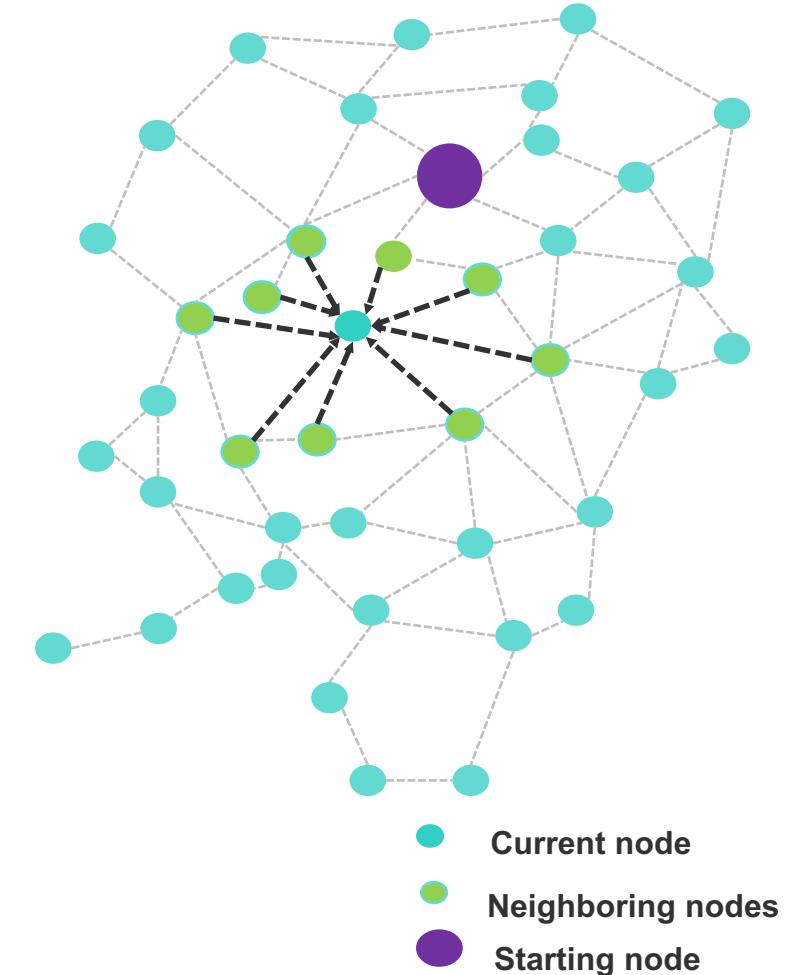
- RNN for 2-D sequences





# RNNs in Various Forms

- RNN for Graph Structures
  - Used in, e.g., image segmentation



[Semantic Object Parsing with Graph LSTM. Liang et al. 2016]





# Outline

- Convolutional Networks (ConvNets)
- Recurrent Networks (RNNs)
  - Long-range dependency, vanishing
  - LSTM
  - RNNs in different forms
- Attention Mechanisms
  - (Query, Key, Value)
  - Attention on Text and Images
- Transformers: Multi-head Attention
  - Transformer
  - BERT





# Attention: Examples

- Chooses which features to pay attention to



A woman is throwing a frisbee in a park.



A dog is standing on a hardwood floor.



A stop sign is on a road with a mountain in the background.



A little girl sitting on a bed with a teddy bear.



A group of people sitting on a boat in the water.



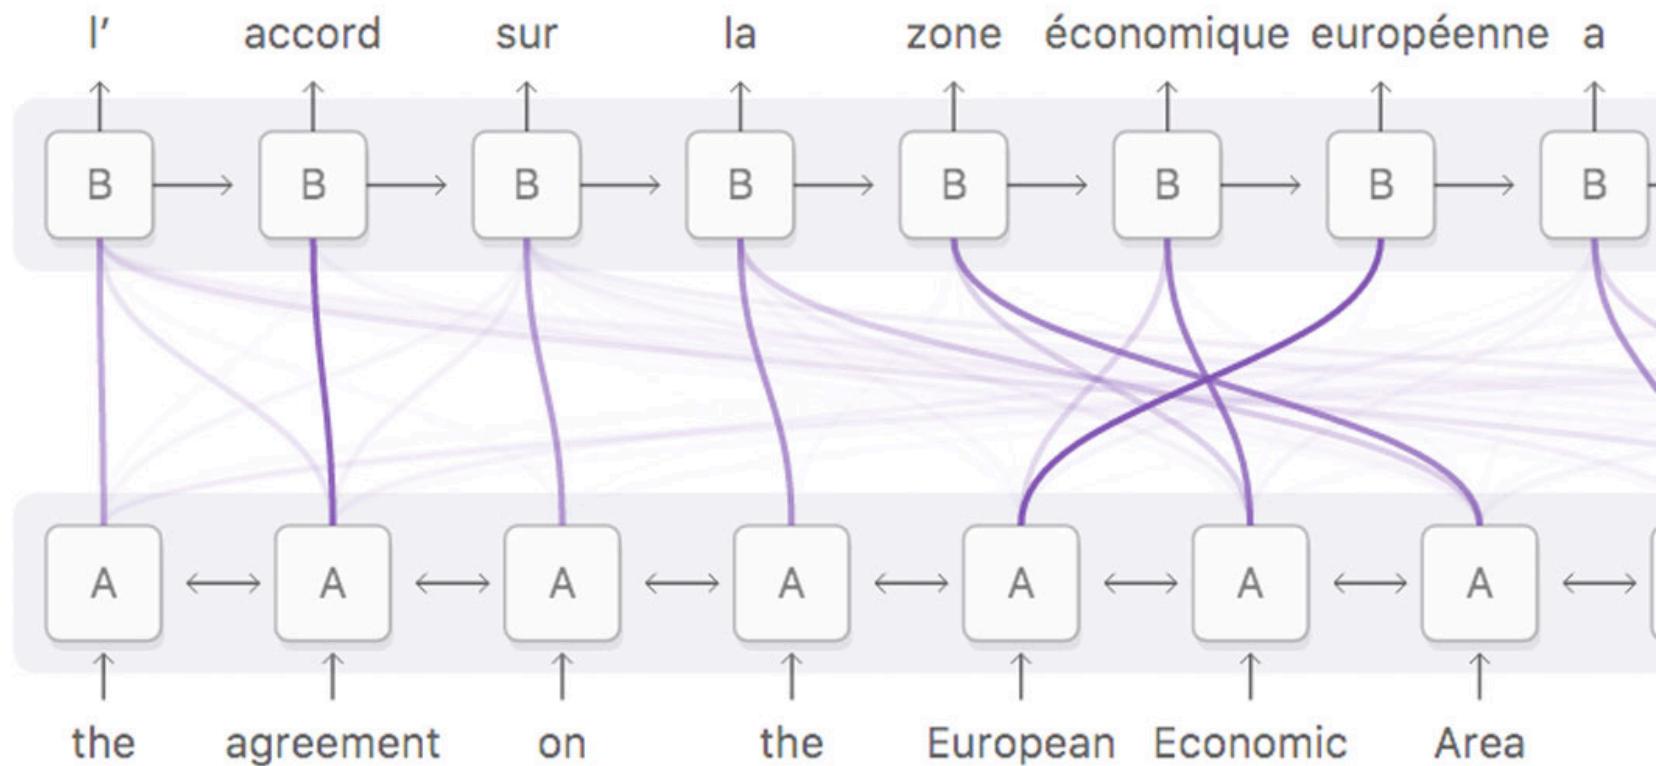
A giraffe standing in a forest with trees in the background.





# Attention: Examples

- Chooses which features to pay attention to



Machine Translation





# Why Attention?





# Why Attention?

- Long-range dependencies
  - Dealing with gradient vanishing problem

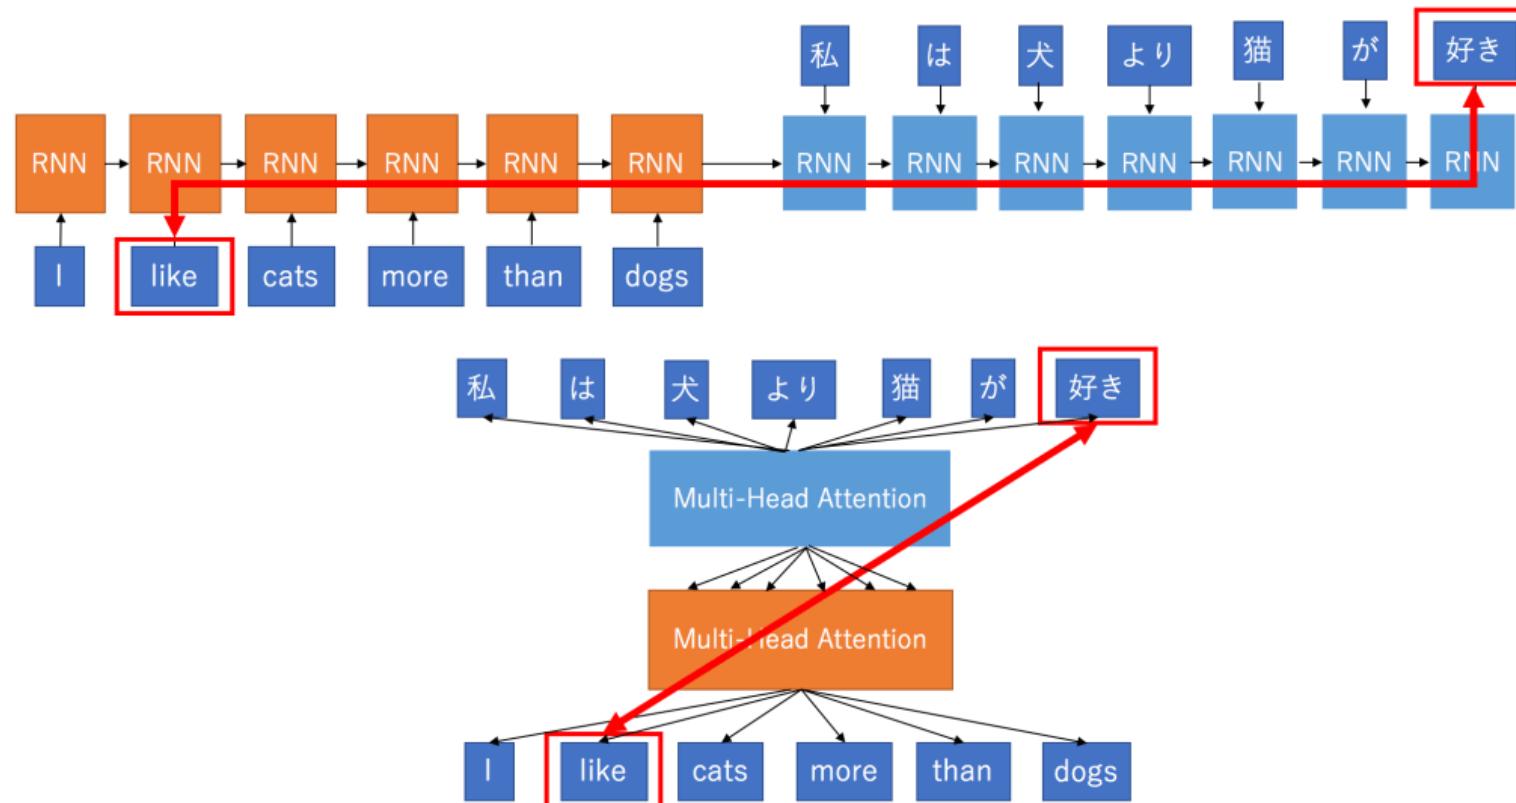


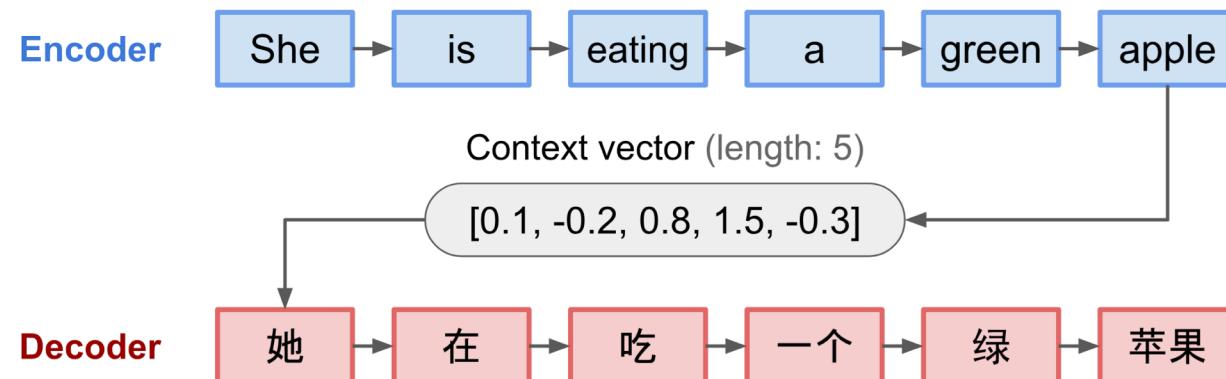
Figure courtesy: [keitakurita](#)





# Why Attention?

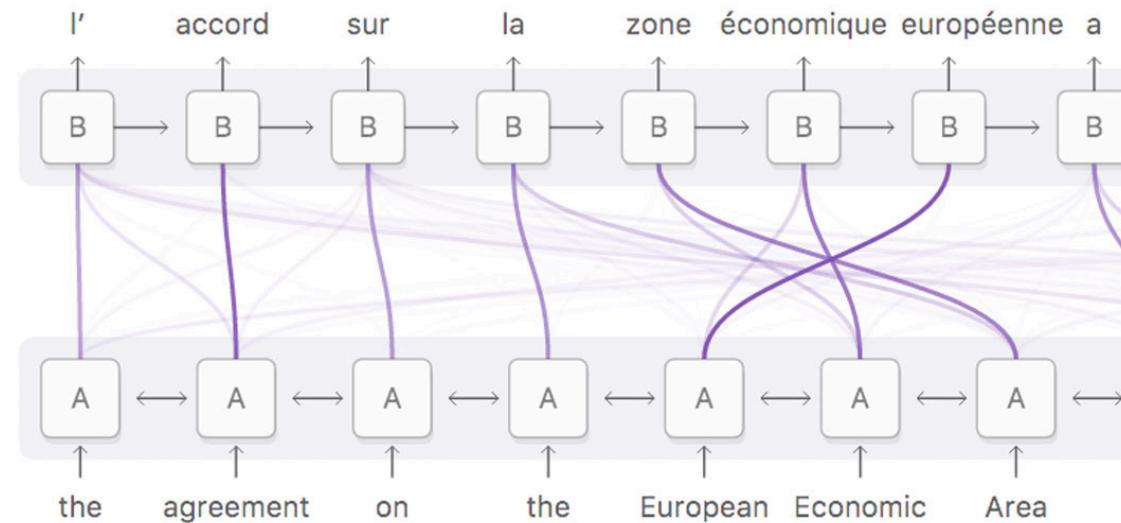
- Long-range dependencies
  - Dealing with gradient vanishing problem
- Fine-grained representation instead of a single global representation
  - Attending to smaller parts of data: patches in images, words in sentences





# Why Attention?

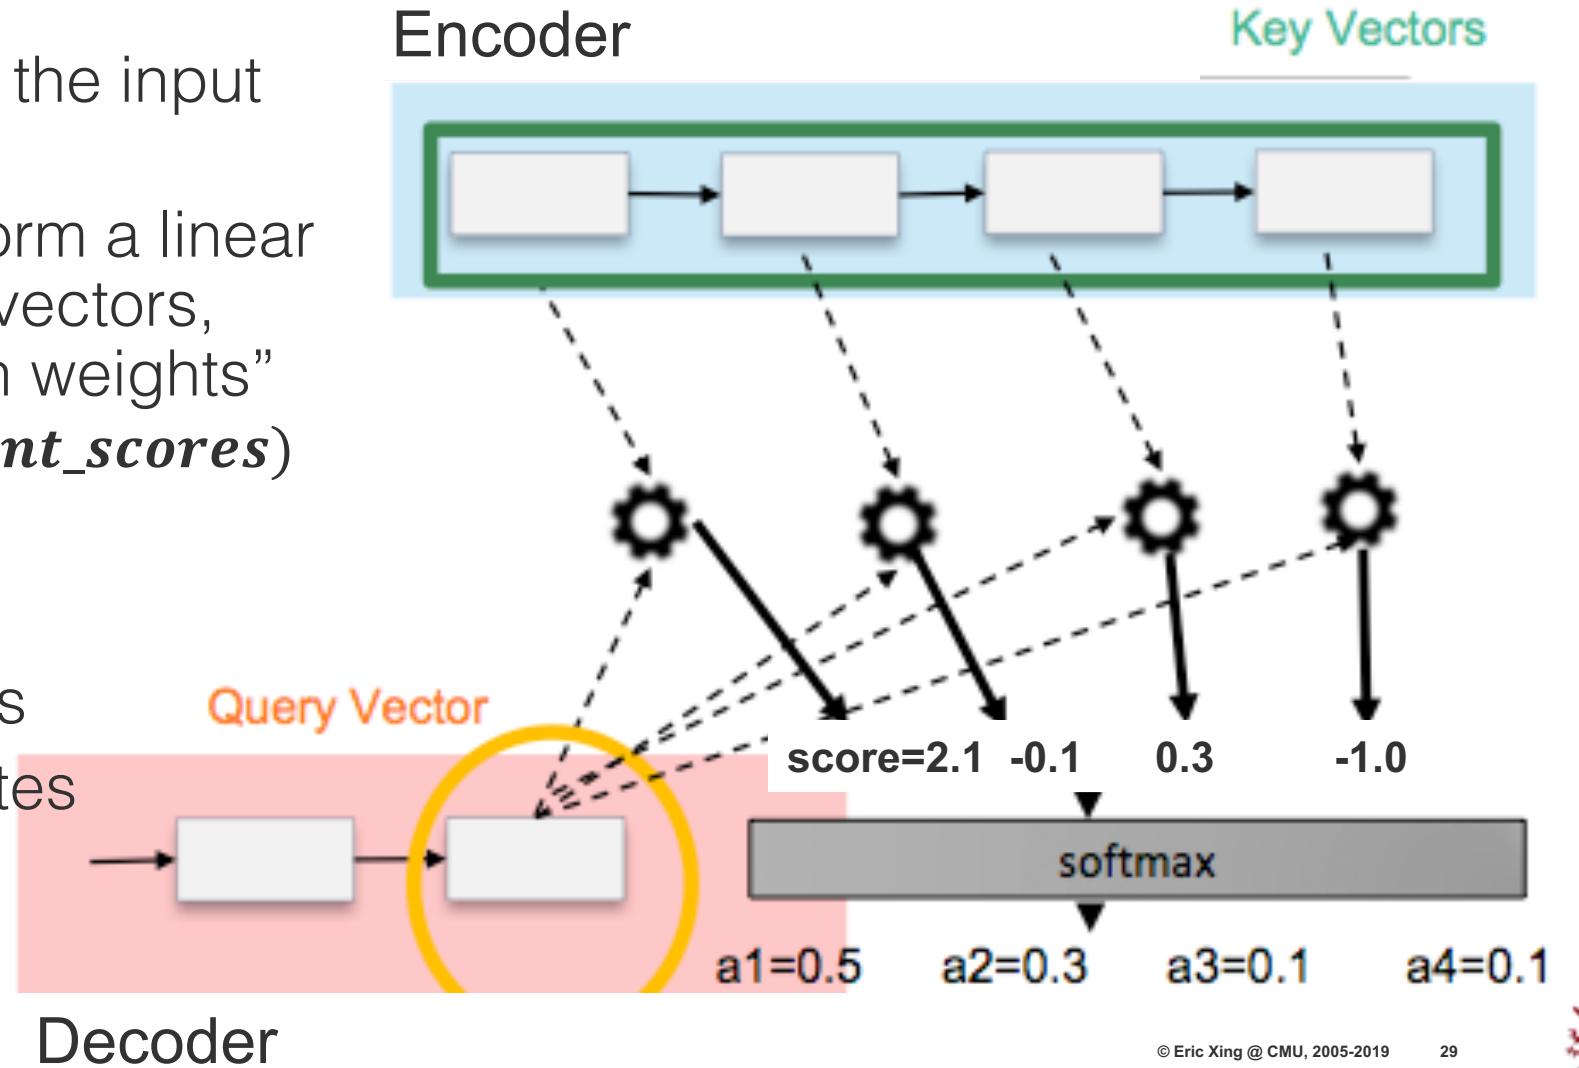
- Long-range dependencies
  - Dealing with gradient vanishing problem
- Fine-grained representation instead of a single global representation
  - Attending to smaller parts of data: patches in images, words in sentences
- Improved Interpretability





# Attention Computation

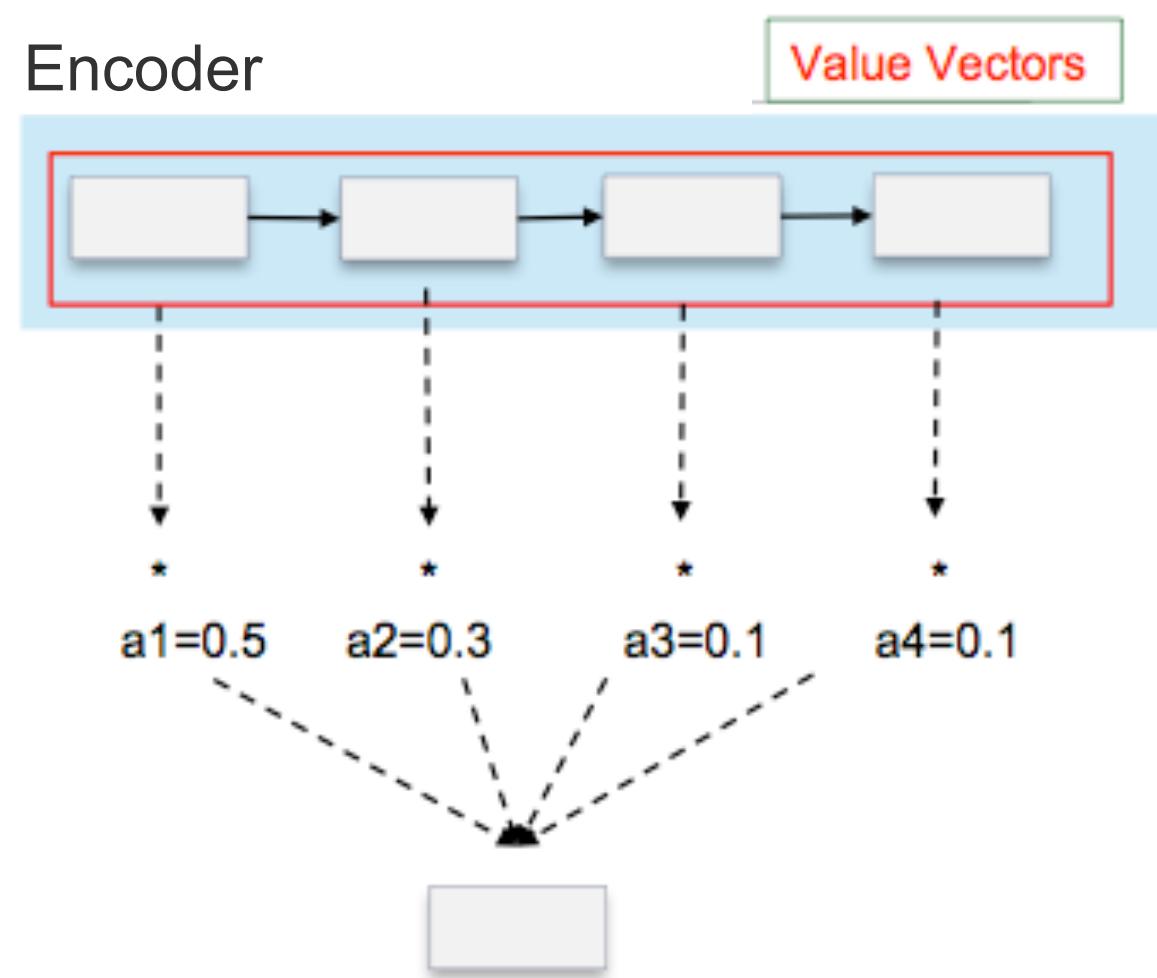
- Encode each token in the input sentence into vectors
- When decoding, perform a linear combination of these vectors, weighted by “attention weights”
  - $a = \text{softmax}(\text{alignment\_scores})$
- Query: decoder state
- Key: all encoder states
- Value: all encoder states





# Attention Computation (cont'd)

- Combine together value by taking the weighted sum
- Query: decoder state
- Key: all encoder states
- Value: all encoder states





# Attention Variants

- Popular attention mechanisms with different alignment score functions

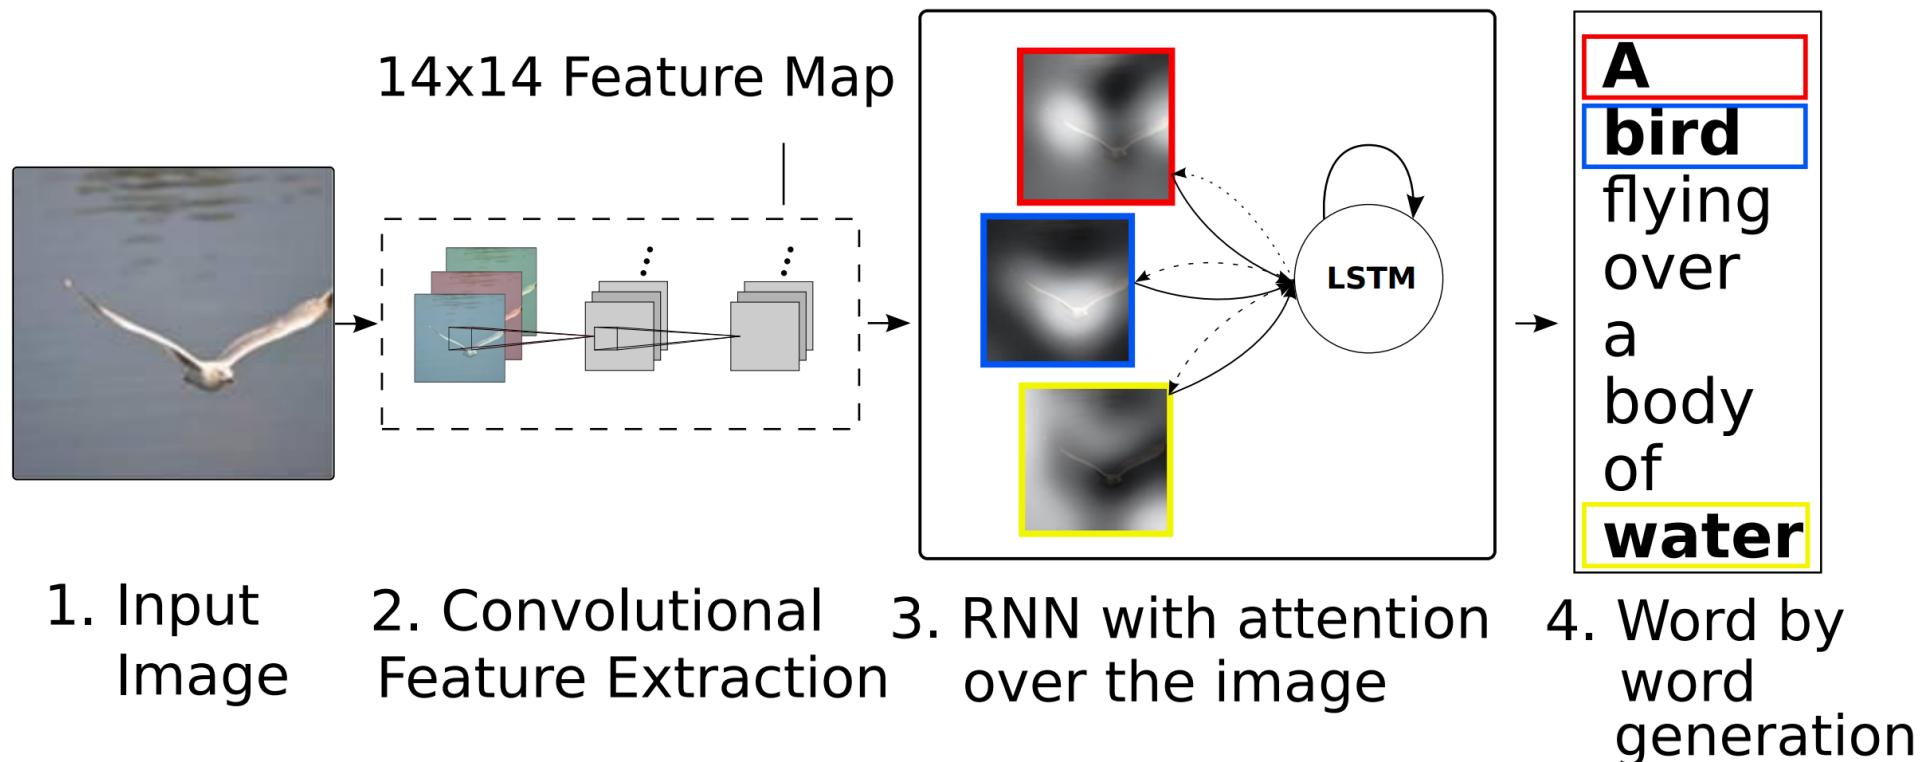
$$\alpha = \text{softmax}(\text{alignment\_scores})$$

- Query:** decoder state  $s_t$
- Key:** all encoder states  $h_i$
- Value:** all encoder states  $h_i$

Name	Alignment score function	Citation
Content-base attention	$\text{score}(s_t, h_i) = \text{cosine}[s_t, h_i]$	Graves2014
Additive(*)	$\text{score}(s_t, h_i) = v_a^\top \tanh(W_a[s_t; h_i])$	Bahdanau2015
Location-Base	$\alpha_{t,i} = \text{softmax}(W_a s_t)$ Note: This simplifies the softmax alignment to only depend on the target position.	Luong2015
General	$\text{score}(s_t, h_i) = s_t^\top W_a h_i$ where $W_a$ is a trainable weight matrix in the attention layer.	Luong2015
Dot-Product	$\text{score}(s_t, h_i) = s_t^\top h_i$	Luong2015
Scaled Dot-Product(^)	$\text{score}(s_t, h_i) = \frac{s_t^\top h_i}{\sqrt{n}}$ Note: very similar to the dot-product attention except for a scaling factor; where n is the dimension of the source hidden state.	Vaswani2017



# Attention on Images – Image Captioning



- **Query:** decoder state
- **Key:** visual feature maps
- **Value:** visual feature maps

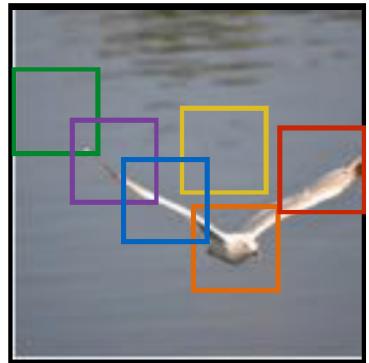




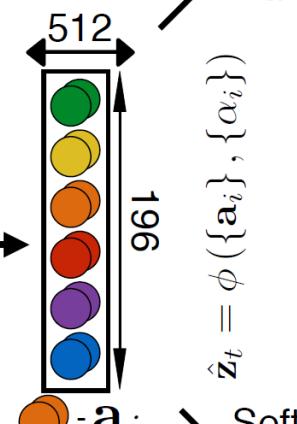
# Attention on Images – Image Captioning

Hard attention vs Soft attention

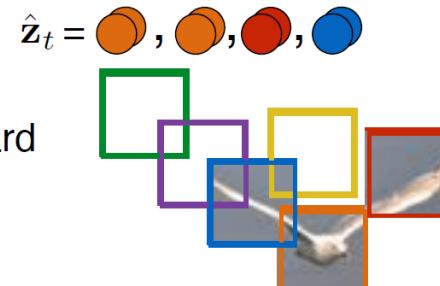
A bird flying over a body of water.



conv-512  
conv-512  
maxpool  
 $14 \times 14 \times 512 = 196 \times 512$  (L x D)  
annotations



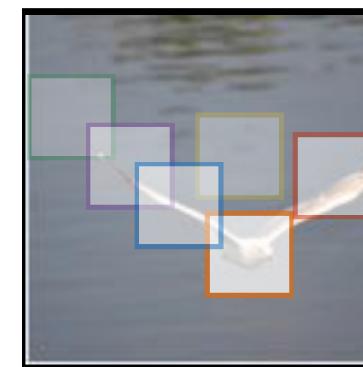
Sample regions of attention



$$L_z = \sum_{z \in \{\text{orange, purple, red, blue}\}} \log p(y | z)$$

$$L_s = \sum_s p(s | \mathbf{a}) \log p(y | s, \mathbf{a})$$

A variational lower bound of maximum likelihood



$$\hat{z}_t = \langle p_1 \ p_2 \ p_3 \ p_4 \ p_5 \ p_6, \ [\text{green circle, yellow circle, orange circle, red circle, purple circle, blue circle}] \rangle$$

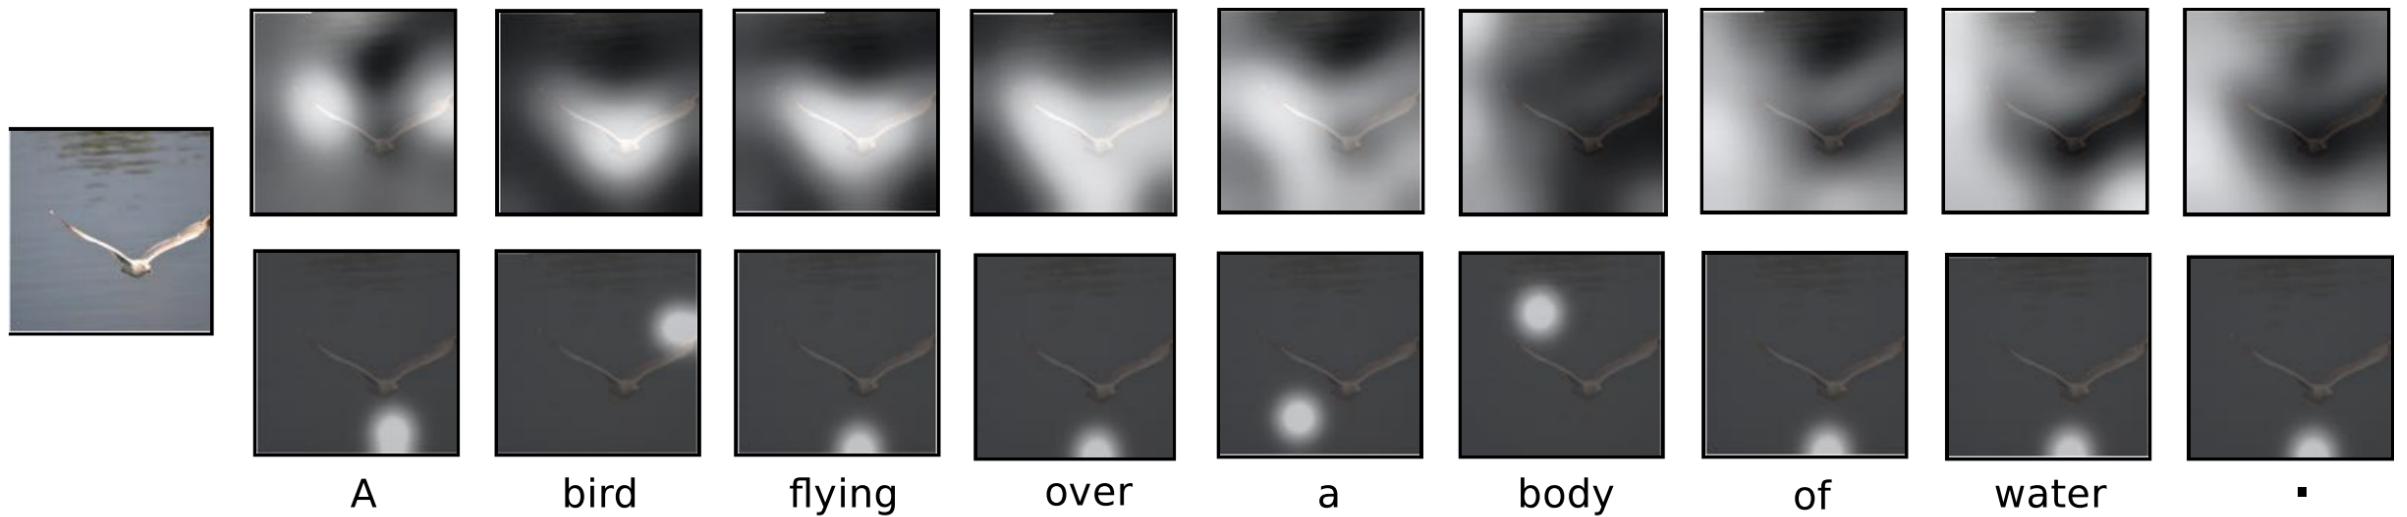
Computes the expected attention





# Attention on Images – Image Captioning

Hard attention vs Soft attention





# Attention on Images – Image Paragraph Generation

- Generate a long paragraph to describe an image
  - Long-term visual and language reasoning
  - Contentful descriptions -- ground sentences on visual features



This picture is taken for three baseball players on a field. The man on the left is wearing a blue baseball cap. The man has a red shirt and white pants. The man in the middle is in a wheelchair and holding a baseball bat. Two men are bending down behind a fence. There are words band on the fence.



A tennis player is attempting to hit the tennis ball with his left foot hand. He is holding a tennis racket. He is wearing a white shirt and white shorts. He has his right arm extended up. There is a crowd of people watching the game. A man is sitting on the chair.

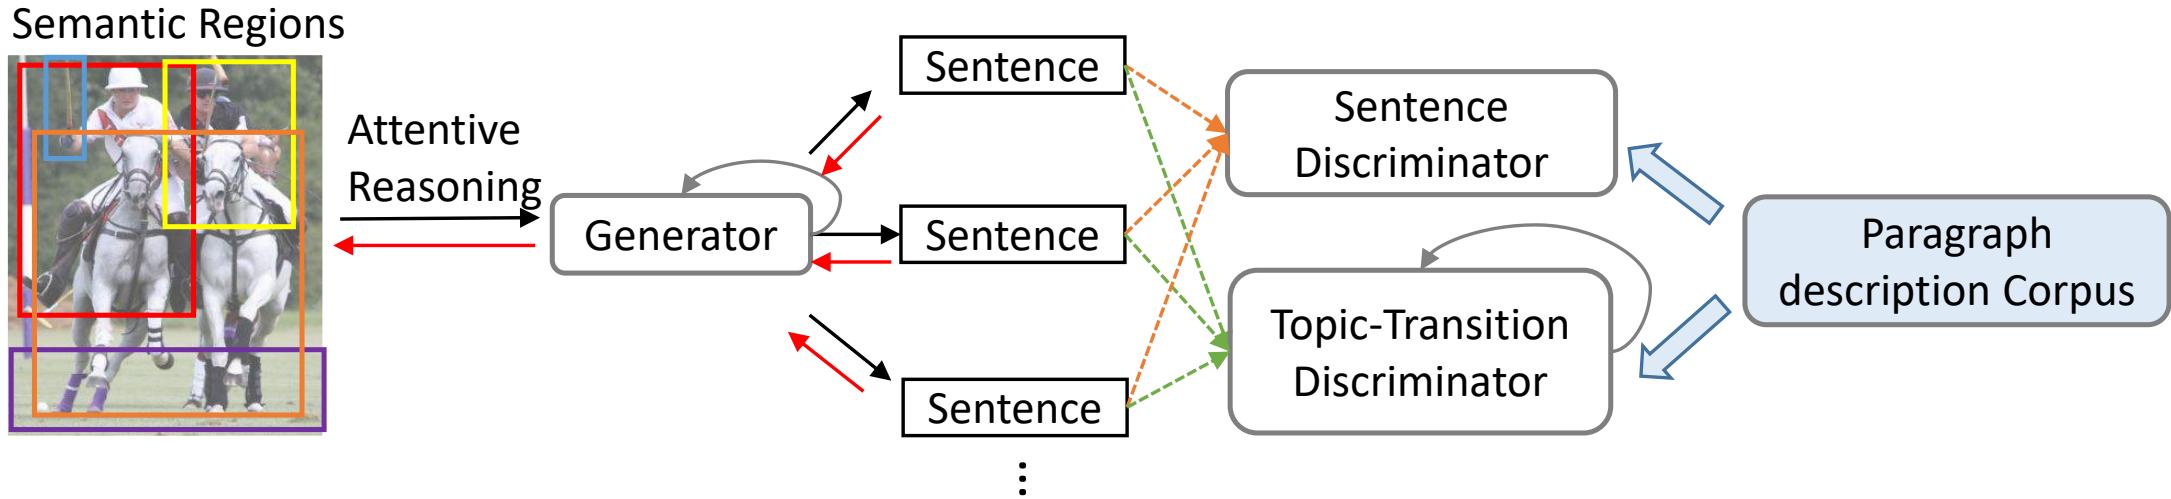


A couple of zebra are standing next to each other on dirt ground near rocks. There are trees behind the zebras. There is a large log on the ground in front of the zebra. There is a large rock formation to the left of the zebra. There is a small hill near a small pond and a wooden log. There are green leaves on the tree.



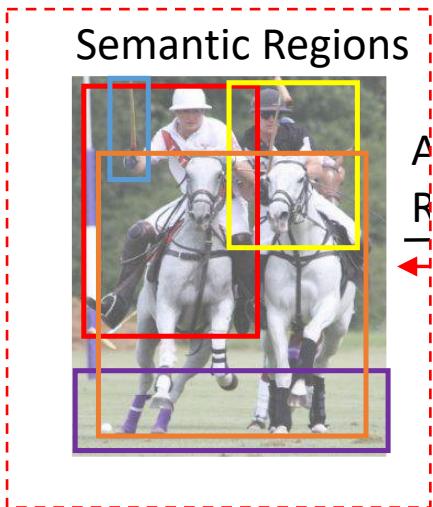


# Attention on Images – Image Paragraph Generation

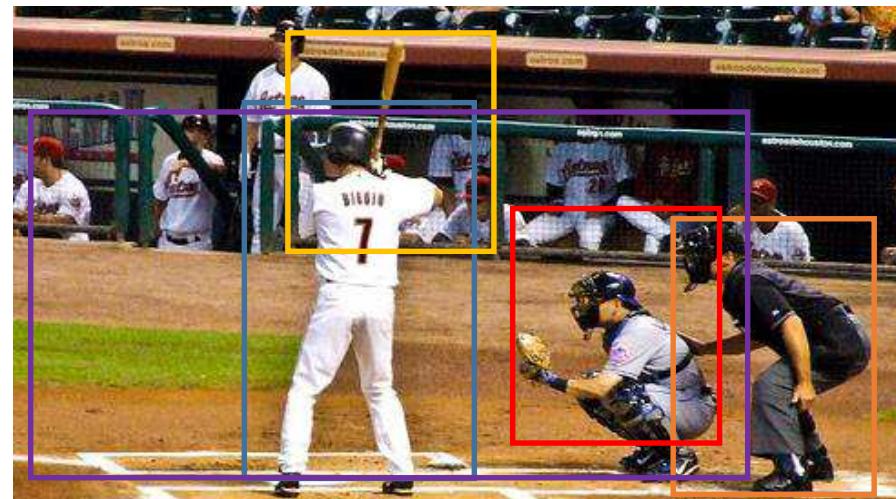




# Attention on Images – Image Paragraph Generation



Semantic region detection & captioning



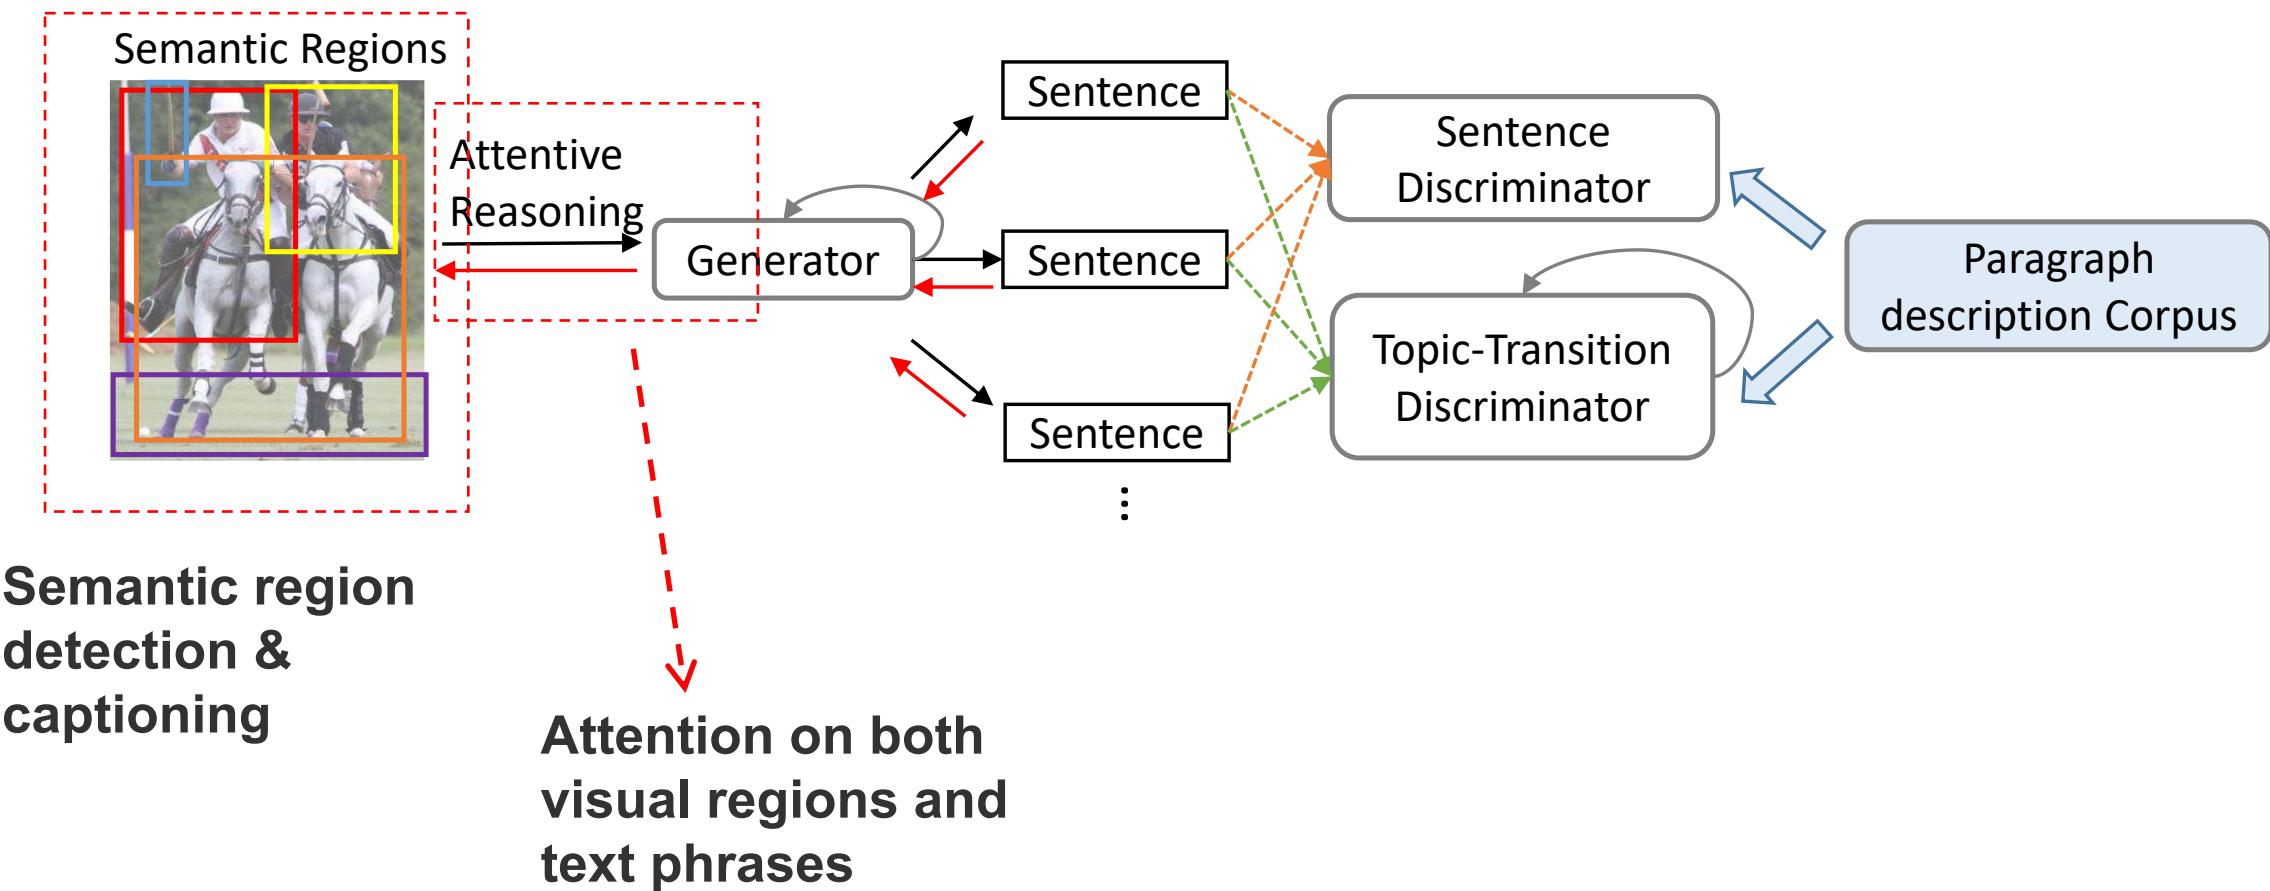
Local Phrases

- people playing baseball
- a man wearing white shirt and pants
- man holding a baseball bat
- person wearing a helmet in the field
- a man bending over



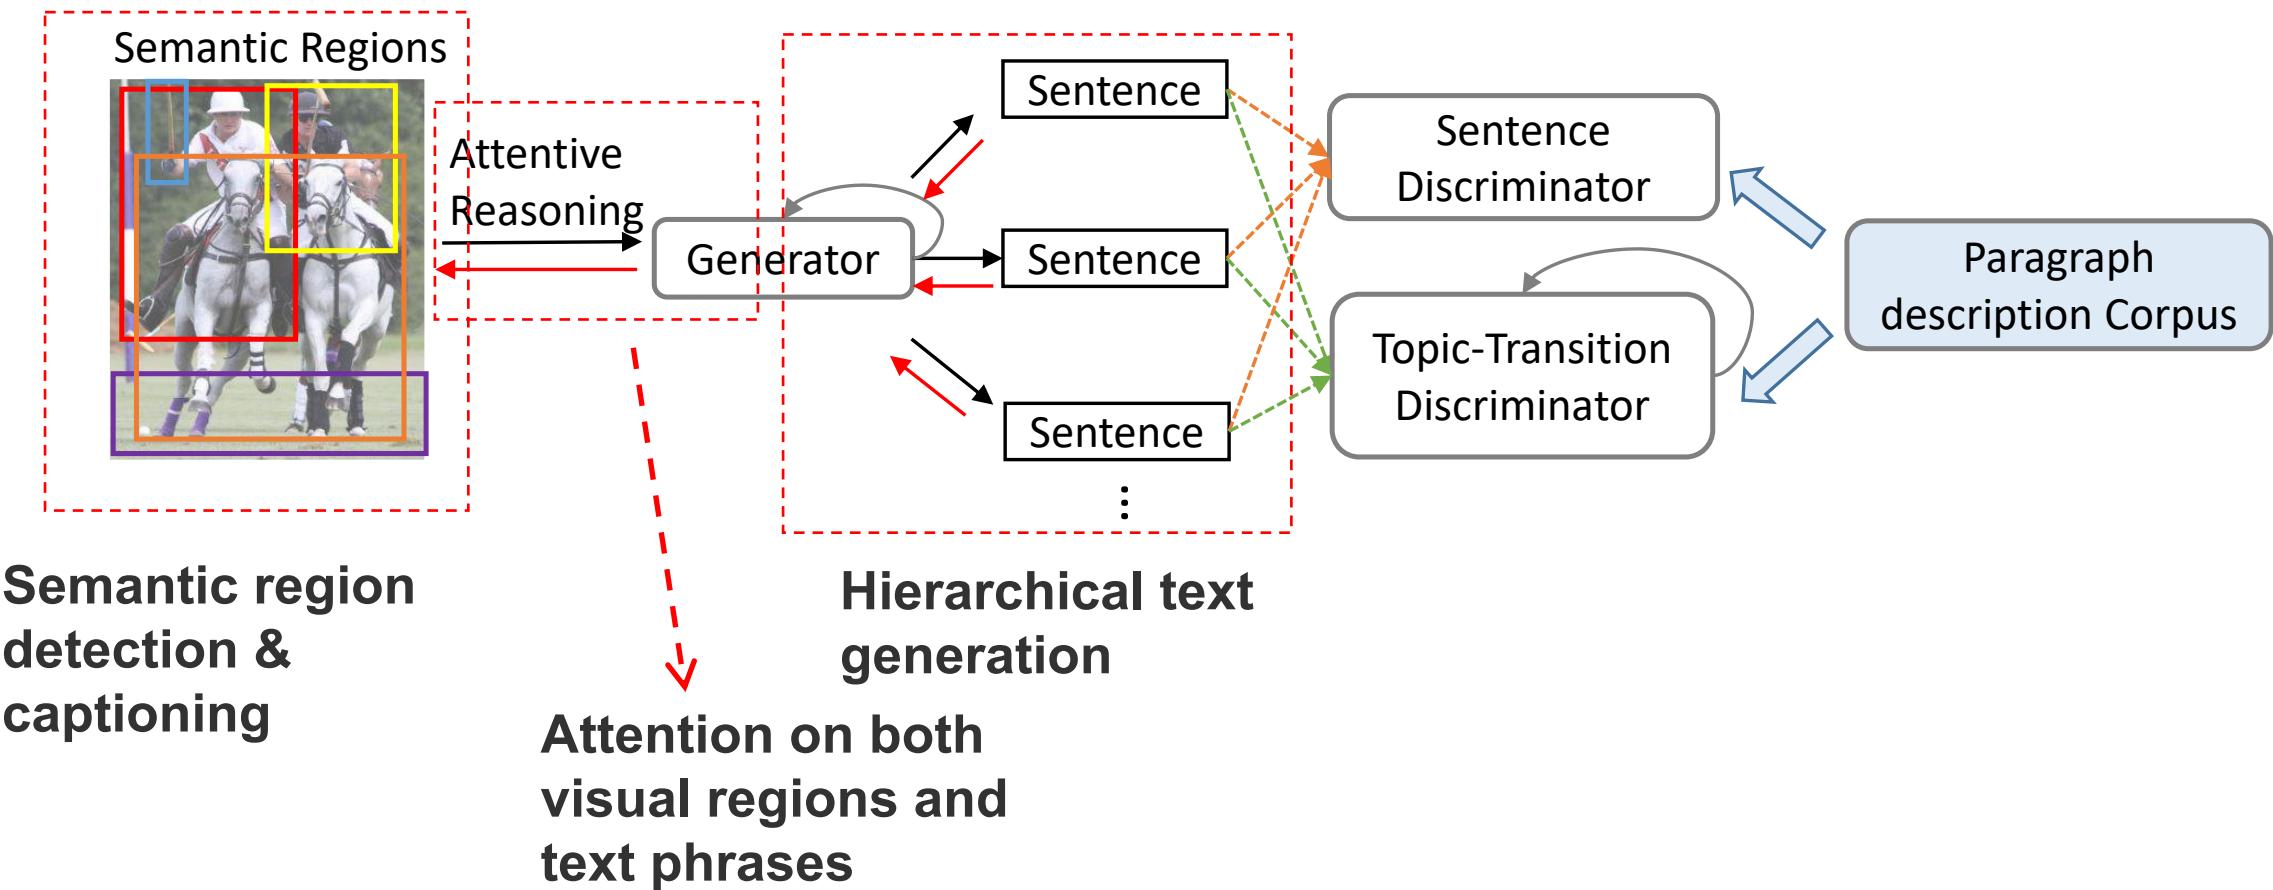


# Attention on Images – Image Paragraph Generation



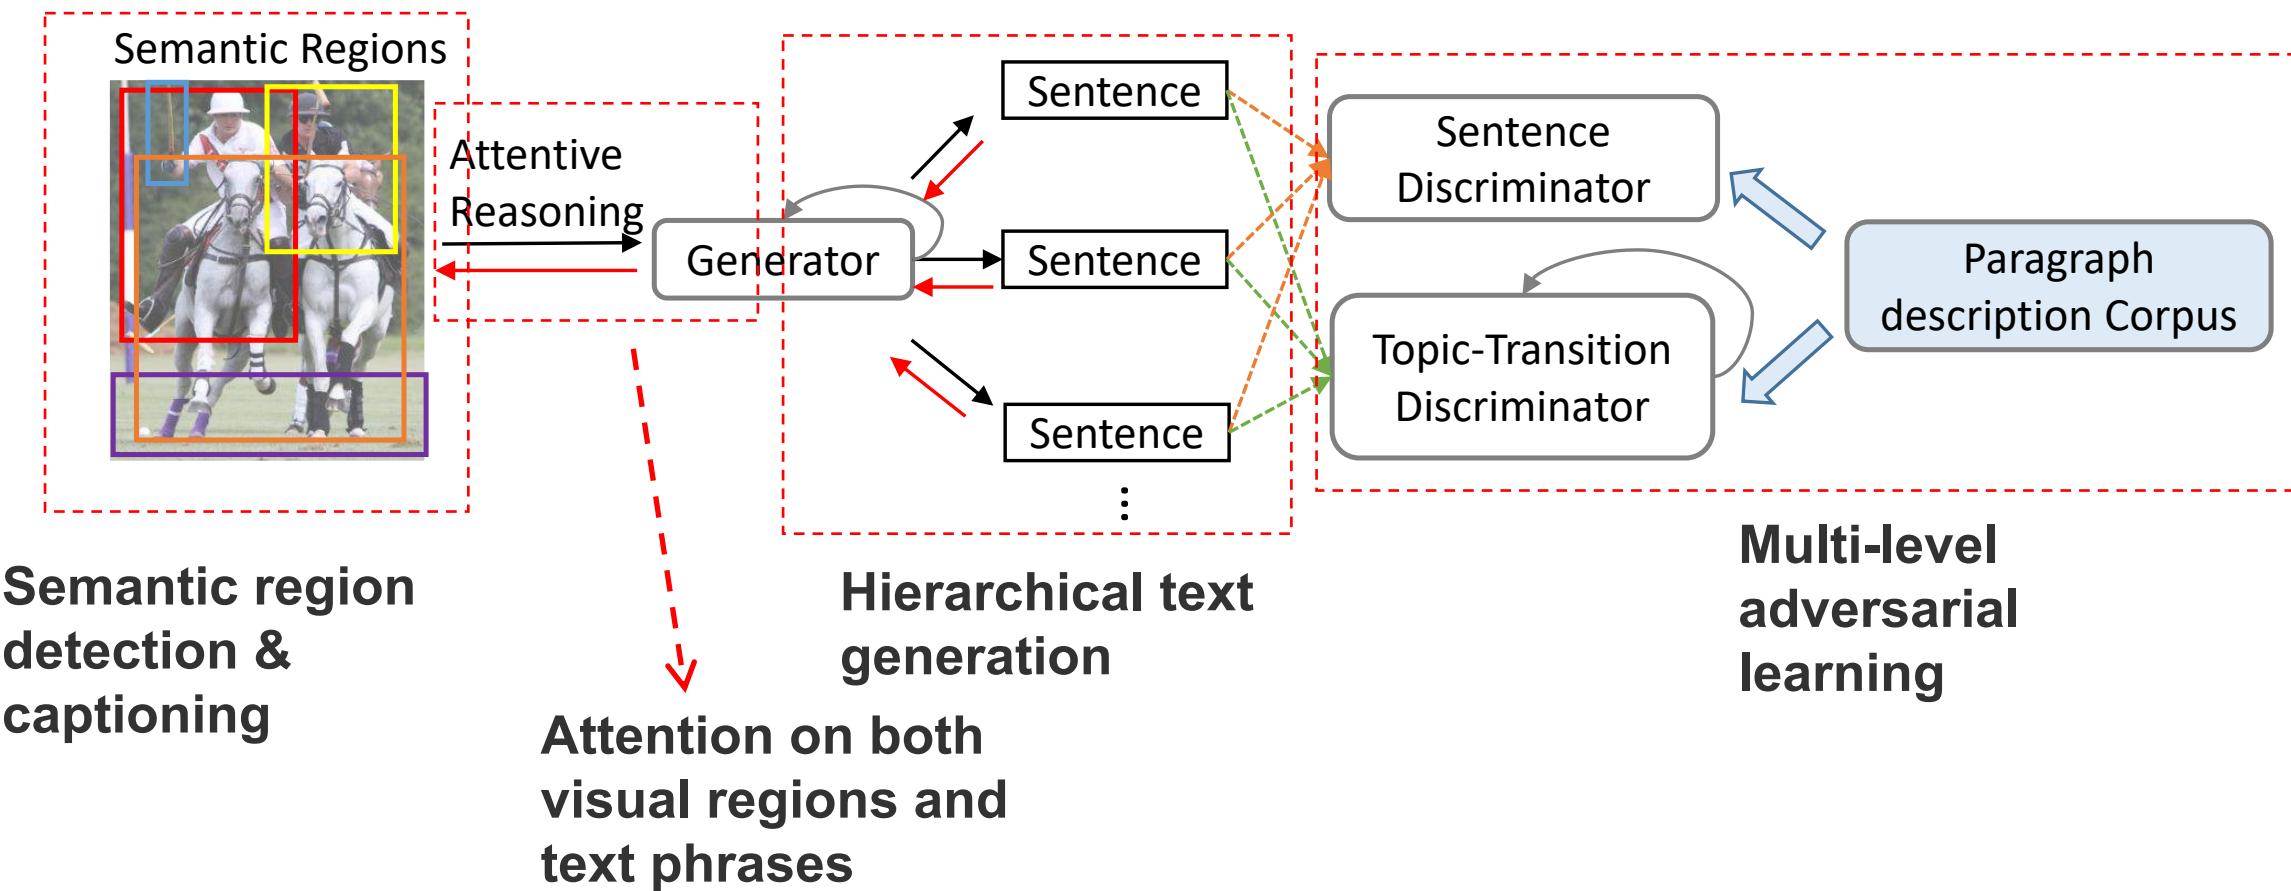


# Attention on Images – Image Paragraph Generation



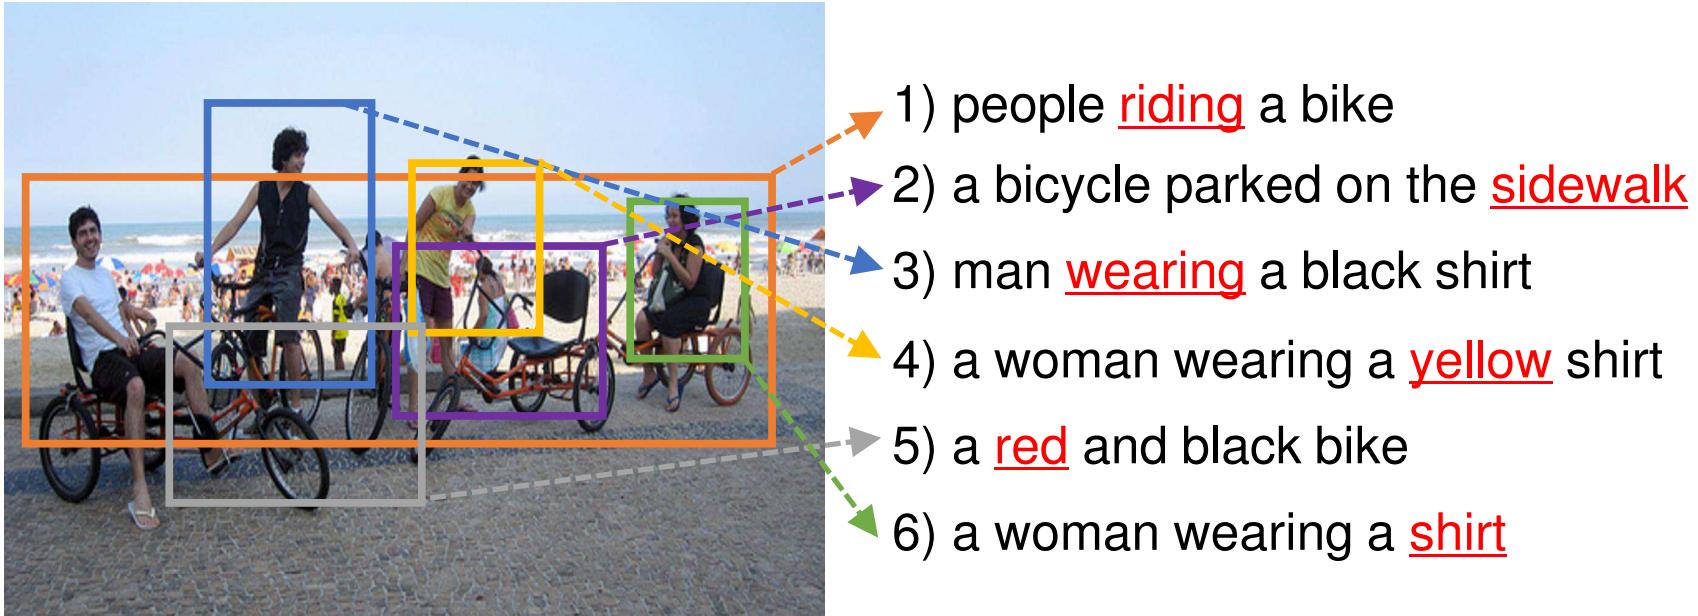


# Attention on Images – Image Paragraph Generation





# Attention on Images – Image Paragraph Generation



**Paragraph:** *A group of people are riding bikes. There are two people riding bikes parked on the sidewalk. He is wearing a black shirt and jeans. A woman is wearing a short sleeve yellow shirt and shorts. There are many other people on the red and black bikes. A woman wearing a shirt is riding a bicycle.*





# Outline

- Convolutional Networks (ConvNets)
- Recurrent Networks (RNNs)
  - Long-range dependency, vanishing
  - LSTM
  - RNNs in different forms
- Attention Mechanisms
  - (Query, Key, Value)
  - Attention on Text and Images
- Transformers: Multi-head Attention
  - Transformer
  - BERT





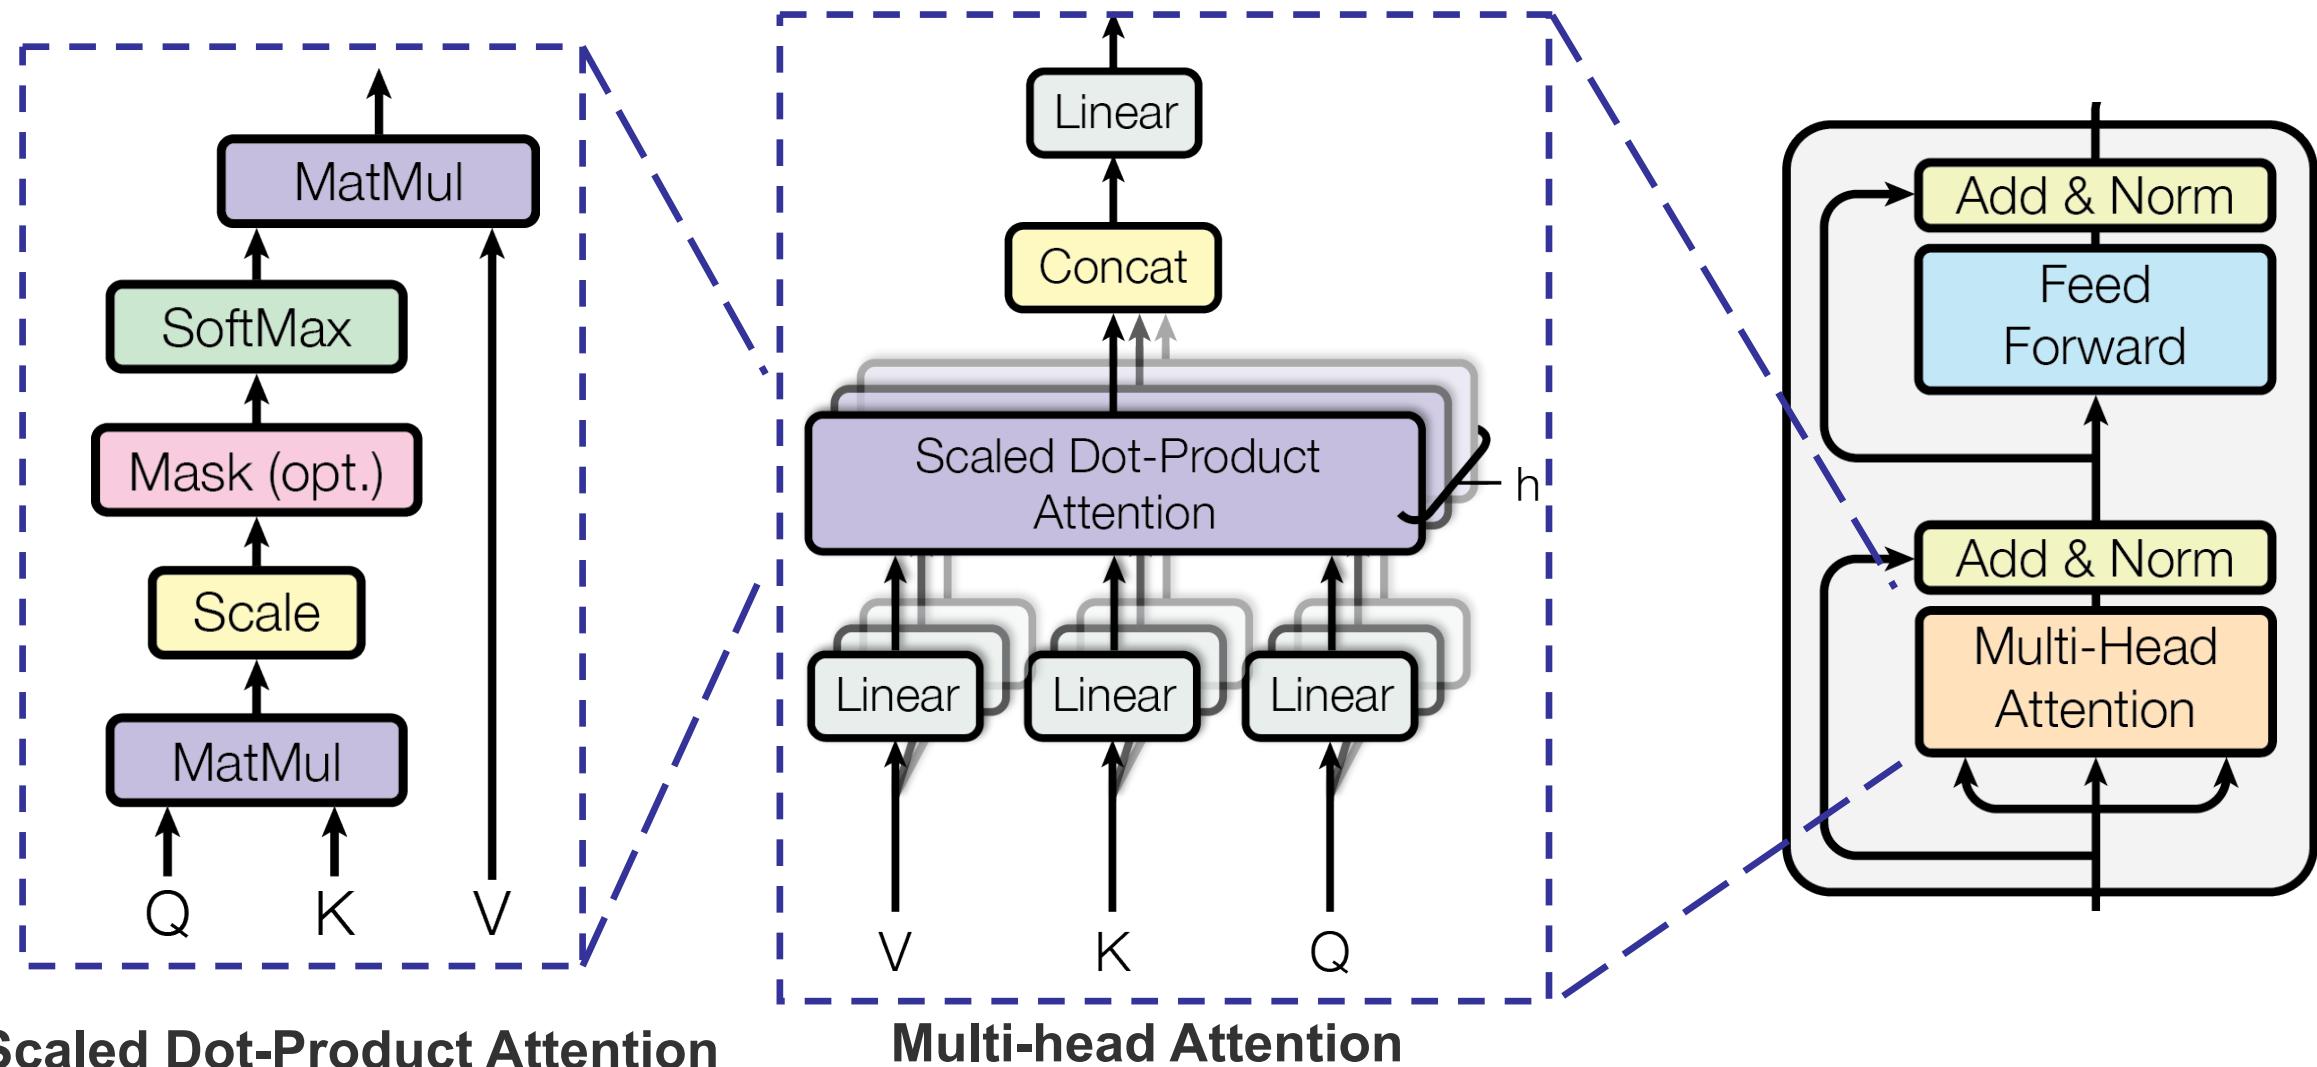
# Transformers – Multi-head (Self-)Attention

- State-of-the-art Results by Transformers
  - [Vaswani et al., 2017] Attention Is All You Need
    - Machine Translation
  - [Devlin et al., 2018] BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding
    - Pre-trained Text Representation
  - [Radford et al., 2019] Language Models are Unsupervised Multitask Learners
    - Language Models





# Multi-head Attention



Scaled Dot-Product Attention

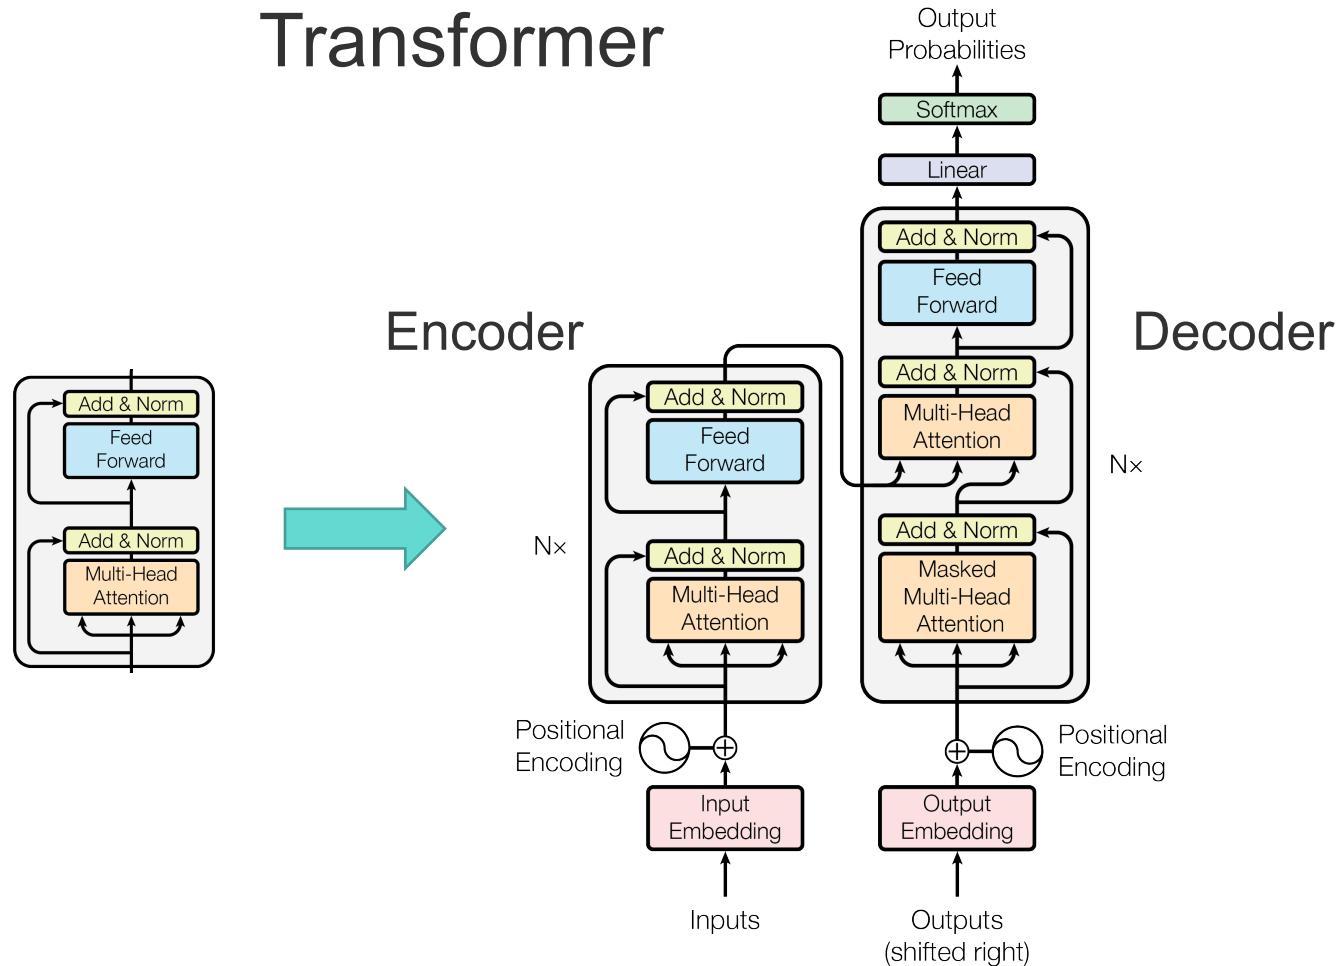
Multi-head Attention





# Multi-head Attention in Encoders and Decoders

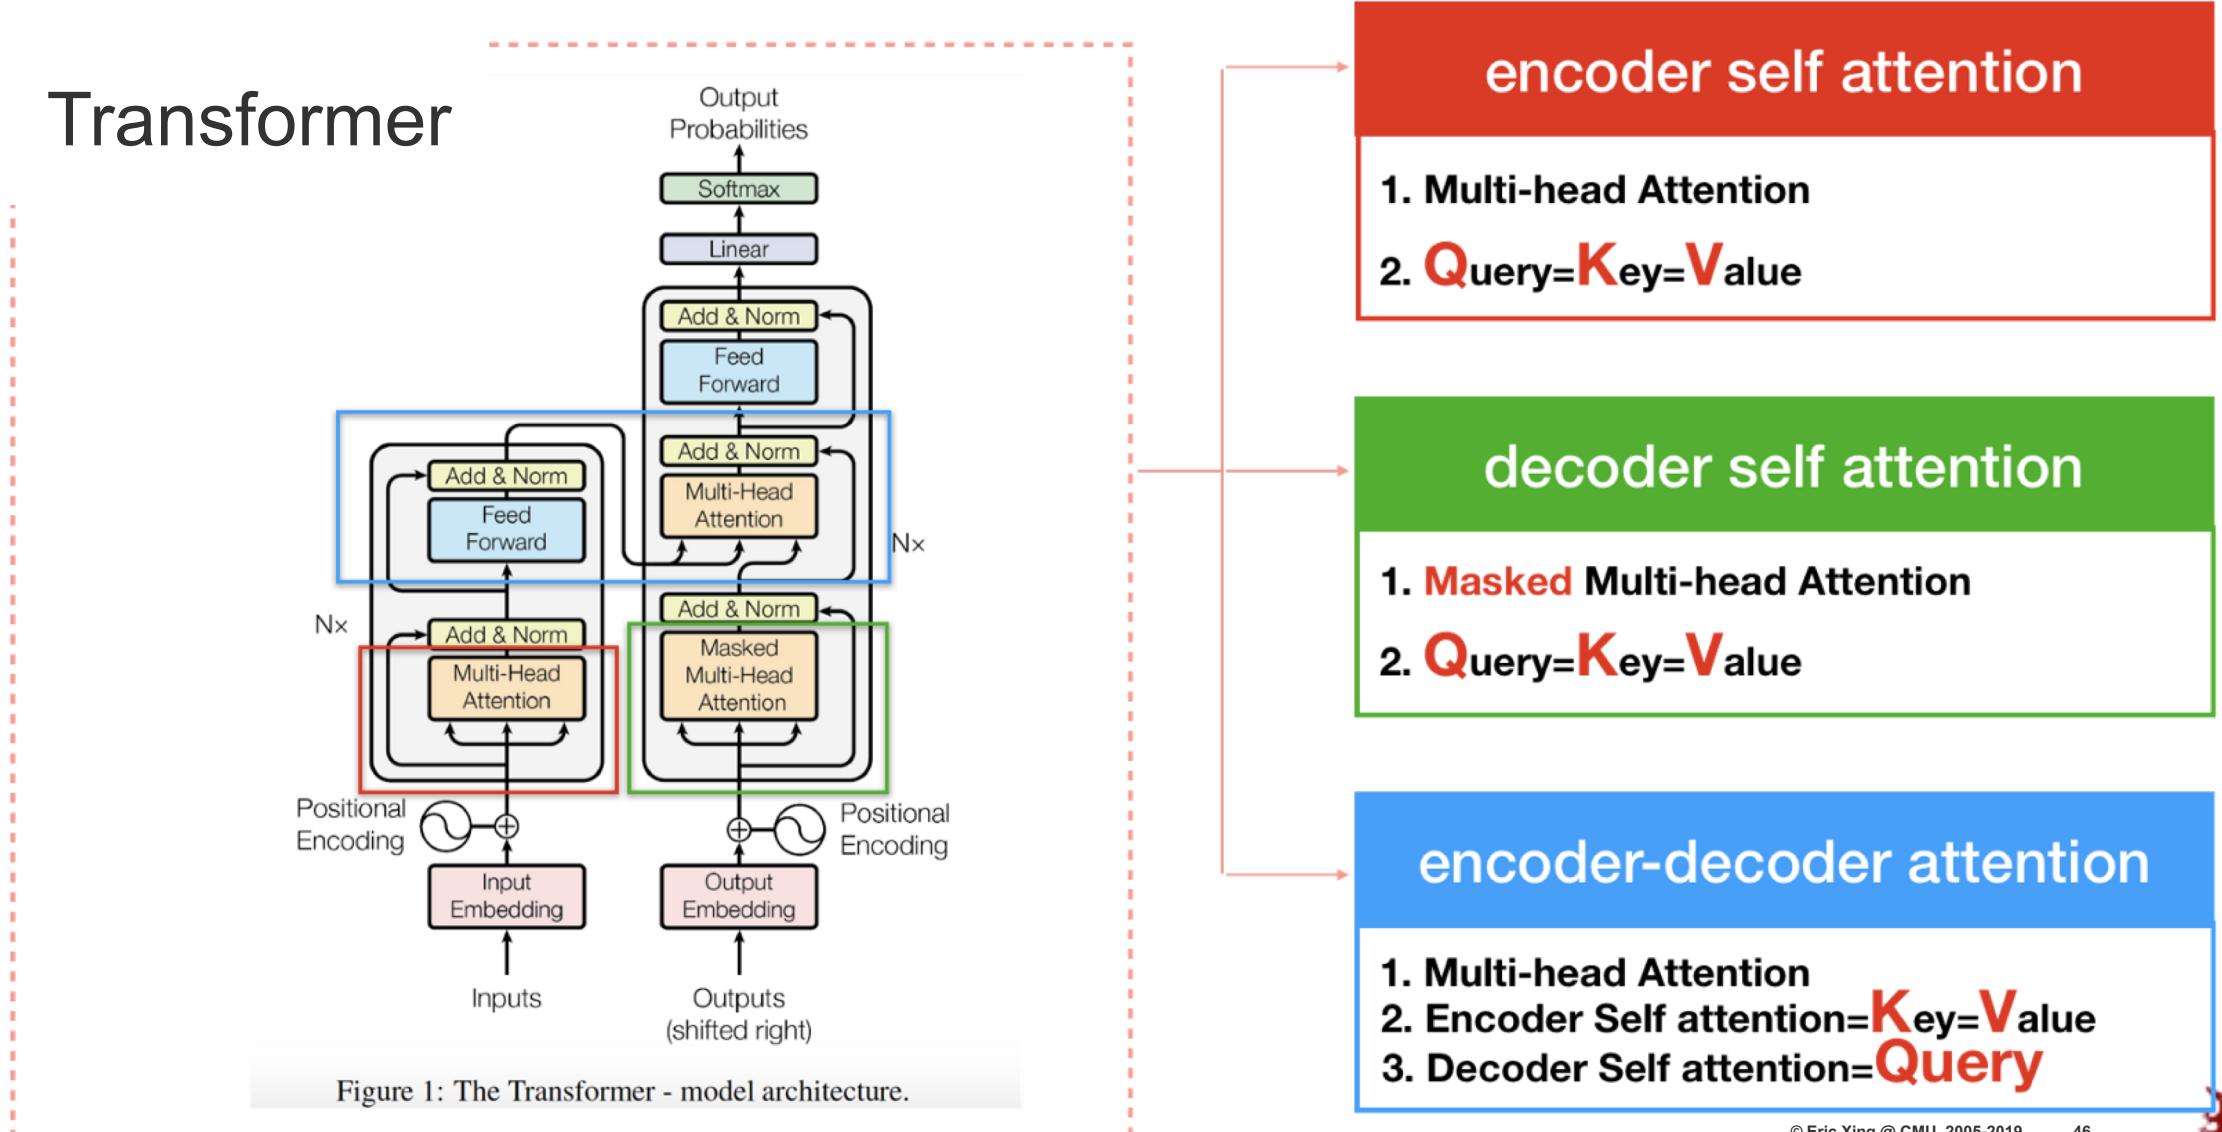
## Transformer





# Multi-head Attention in Encoders and Decoders

## Transformer





# BERT: Pre-trained Text Representation Model

- Conventional word embedding:
  - Word2vec, Glove
  - A pre-trained matrix, each row is an embedding vector of a word

	0	1	2	3	4	5	6	7	8	9	...
<b>fox</b>	-0.348680	-0.077720	0.177750	-0.094953	-0.452890	0.237790	0.209440	0.037886	0.035064	0.899010	...
<b>ham</b>	-0.773320	-0.282540	0.580760	0.841480	0.258540	0.585210	-0.021890	-0.463680	0.139070	0.658720	...
<b>brown</b>	-0.374120	-0.076264	0.109260	0.186620	0.029943	0.182700	-0.631980	0.133060	-0.128980	0.603430	...
<b>beautiful</b>	0.171200	0.534390	-0.348540	-0.097234	0.101800	-0.170860	0.295650	-0.041816	-0.516550	2.117200	...
<b>jumps</b>	-0.334840	0.215990	-0.350440	-0.260020	0.411070	0.154010	-0.386110	0.206380	0.386700	1.460500	...
<b>eggs</b>	-0.417810	-0.035192	-0.126150	-0.215930	-0.669740	0.513250	-0.797090	-0.068611	0.634660	1.256300	...
<b>beans</b>	-0.423290	-0.264500	0.200870	0.082187	0.066944	1.027600	-0.989140	-0.259950	0.145960	0.766450	...
<b>sky</b>	0.312550	-0.303080	0.019587	-0.354940	0.100180	-0.141530	-0.514270	0.886110	-0.530540	1.556600	...
<b>bacon</b>	-0.430730	-0.016025	0.484620	0.101390	-0.299200	0.761820	-0.353130	-0.325290	0.156730	0.873210	...
<b>breakfast</b>	0.073378	0.227670	0.208420	-0.456790	-0.078219	0.601960	-0.024494	-0.467980	0.054627	2.283700	...
<b>toast</b>	0.130740	-0.193730	0.253270	0.090102	-0.272580	-0.030571	0.096945	-0.115060	0.484000	0.848380	...
<b>today</b>	-0.156570	0.594890	-0.031445	-0.077586	0.278630	-0.509210	-0.066350	-0.081890	-0.047986	2.803600	...
<b>blue</b>	0.129450	0.036518	0.032298	-0.060034	0.399840	-0.103020	-0.507880	0.076630	-0.422920	0.815730	...
<b>green</b>	-0.072368	0.233200	0.137260	-0.156630	0.248440	0.349870	-0.241700	-0.091426	-0.530150	1.341300	...
<b>kings</b>	0.259230	-0.854690	0.360010	-0.642000	0.568530	-0.321420	0.173250	0.133030	-0.089720	1.528600	...
<b>dog</b>	-0.057120	0.052685	0.003026	-0.048517	0.007043	0.041856	-0.024704	-0.039783	0.009614	0.308416	...
<b>sausages</b>	-0.174290	-0.064869	-0.046976	0.287420	-0.128150	0.647630	0.056315	-0.240440	-0.025094	0.502220	...
<b>lazy</b>	-0.353320	-0.299710	-0.176230	-0.321940	-0.385640	0.586110	0.411160	-0.418680	0.073093	1.486500	...
<b>love</b>	0.139490	0.534530	-0.252470	-0.125650	0.048748	0.152440	0.199060	-0.065970	0.128830	2.055900	...
<b>quick</b>	-0.445630	0.191510	-0.249210	0.465900	0.161950	0.212780	-0.046480	0.021170	0.417660	1.686900	...

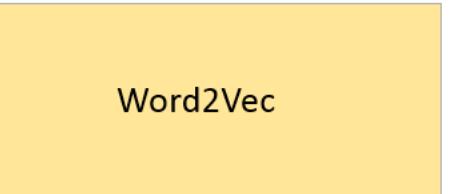


# BERT: Pre-trained Text Representation Model

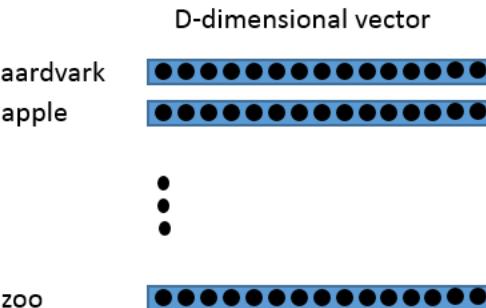
- Conventional word embedding:
  - Word2vec, Glove
  - A pre-trained matrix, each row is an embedding vector of a word

English Wikipedia Corpus

The Annual Reminder continued through July 4, 1969. This final Annual Reminder took place less than a week after the June 28 Stonewall riots, in which the patrons of the Stonewall Inn, a gay bar in Greenwich Village, fought against police who raided the bar. Rodwell received several telephone calls threatening him and the other New York participants, but he was able to arrange for police protection for the chartered bus all the way to Philadelphia. About 45 people participated, including the deputy mayor of Philadelphia and his wife. The dress code was still in effect at the Reminder, but two women from the New York contingent broke from the single-file picket line and held hands. When Kameny tried to break them apart, Rodwell furiously denounced him to onlooking members of the press. Following the 1969 Annual Reminder, there was a sense, particularly among the younger and more radical participants, that the time for silent picketing had passed. Dissent and dissatisfaction had begun to take new and more emphatic forms in society.<sup>10</sup> The conference passed a resolution drafted by Rodwell, his partner Fred Sargent, Brody and Linda Rhodes to move the demonstration from July 4 in Philadelphia to the last weekend in June in New York City, as well as proposing to "other organizations throughout the country... suggesting that they hold parallel demonstrations on that day" to commemorate the Stonewall riot. ....



Embedding Matrix



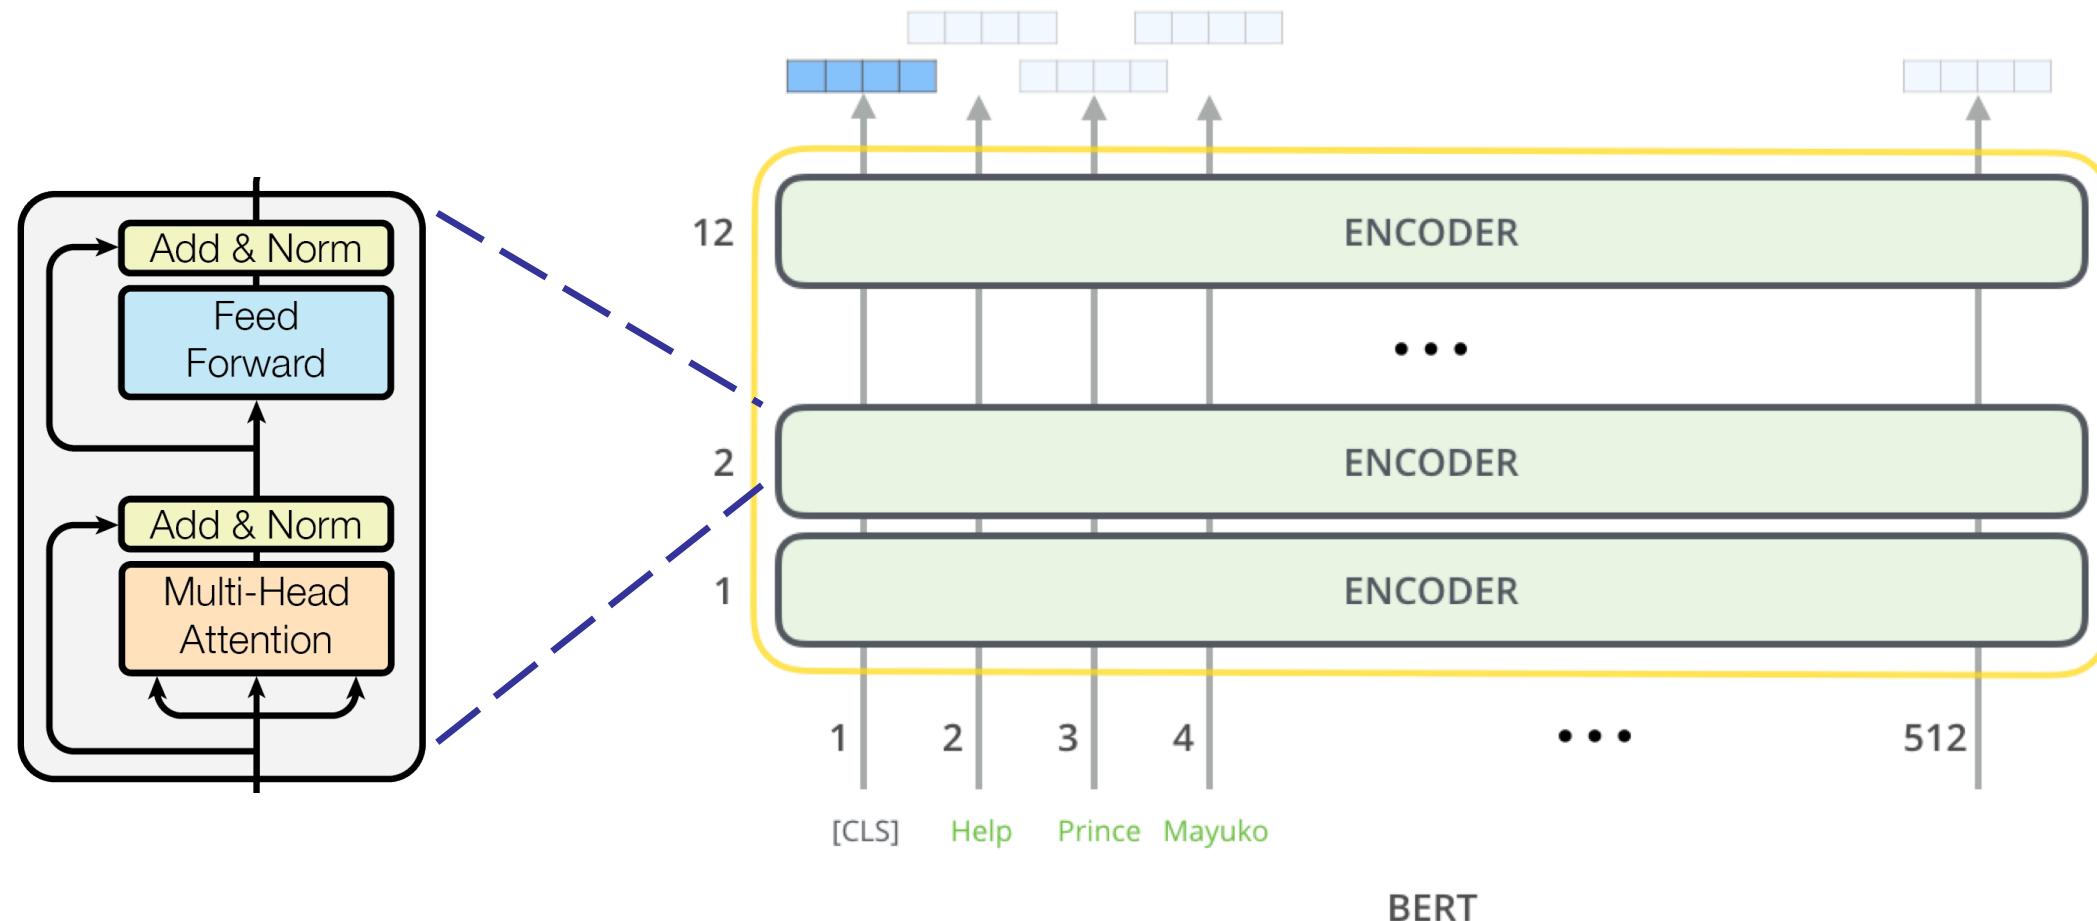
Word2Vec

	0	1	2	3	4	5	6	7	8	9	...
<b>fox</b>	-0.348680	-0.077720	0.177750	-0.094953	-0.452890	0.237790	0.209440	0.037886	0.035064	0.899010	...
<b>ham</b>	-0.773320	-0.282540	0.580760	0.841480	0.258540	0.585210	-0.021890	-0.463680	0.139070	0.658720	...
<b>brown</b>	-0.374120	-0.076264	0.109260	0.186620	0.029943	0.182700	-0.631980	0.133060	-0.128980	0.603430	...
<b>beautiful</b>	0.171200	0.534390	-0.348540	-0.097234	0.101800	-0.170860	0.295650	-0.041816	-0.516550	2.117200	...
<b>jumps</b>	-0.334840	0.215990	-0.350440	-0.260020	0.411070	0.154010	-0.386110	0.206380	0.386700	1.460500	...
<b>eggs</b>	-0.417810	-0.035192	-0.126150	-0.215930	-0.669740	0.513250	-0.797090	-0.068611	0.634660	1.256300	...
<b>beans</b>	-0.423290	-0.264500	0.200870	0.082187	0.066944	1.027600	-0.989140	-0.259950	0.145960	0.766450	...
<b>sky</b>	0.312550	-0.303080	0.019587	-0.354940	0.100180	-0.141530	-0.514270	0.886110	-0.530540	1.556600	...
<b>bacon</b>	-0.430730	-0.016025	0.484620	0.101390	-0.299200	0.761820	-0.353130	-0.325290	0.156730	0.873210	...
<b>breakfast</b>	0.073378	0.227670	0.208420	-0.456790	-0.078219	0.601960	-0.024494	-0.467980	0.054627	2.283700	...
<b>...</b>	<b>...</b>	<b>...</b>	<b>...</b>	<b>...</b>	<b>...</b>	<b>...</b>	<b>...</b>	<b>...</b>	<b>...</b>	<b>...</b>	...



# BERT: Pre-trained Text Representation Model

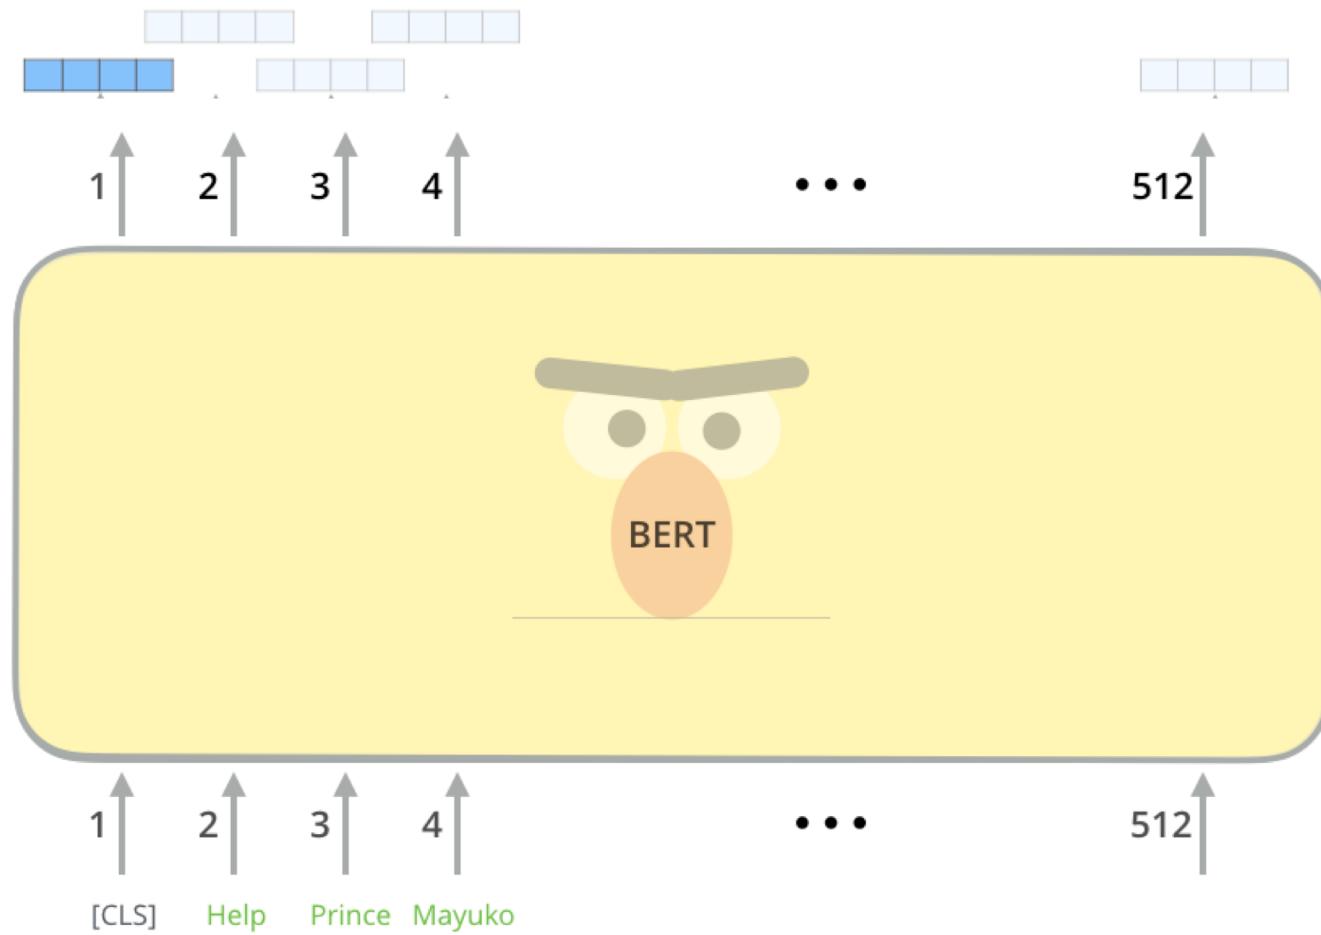
- A **model** to extract *contextualized* word embedding





# BERT: Pre-trained Text Representation Model

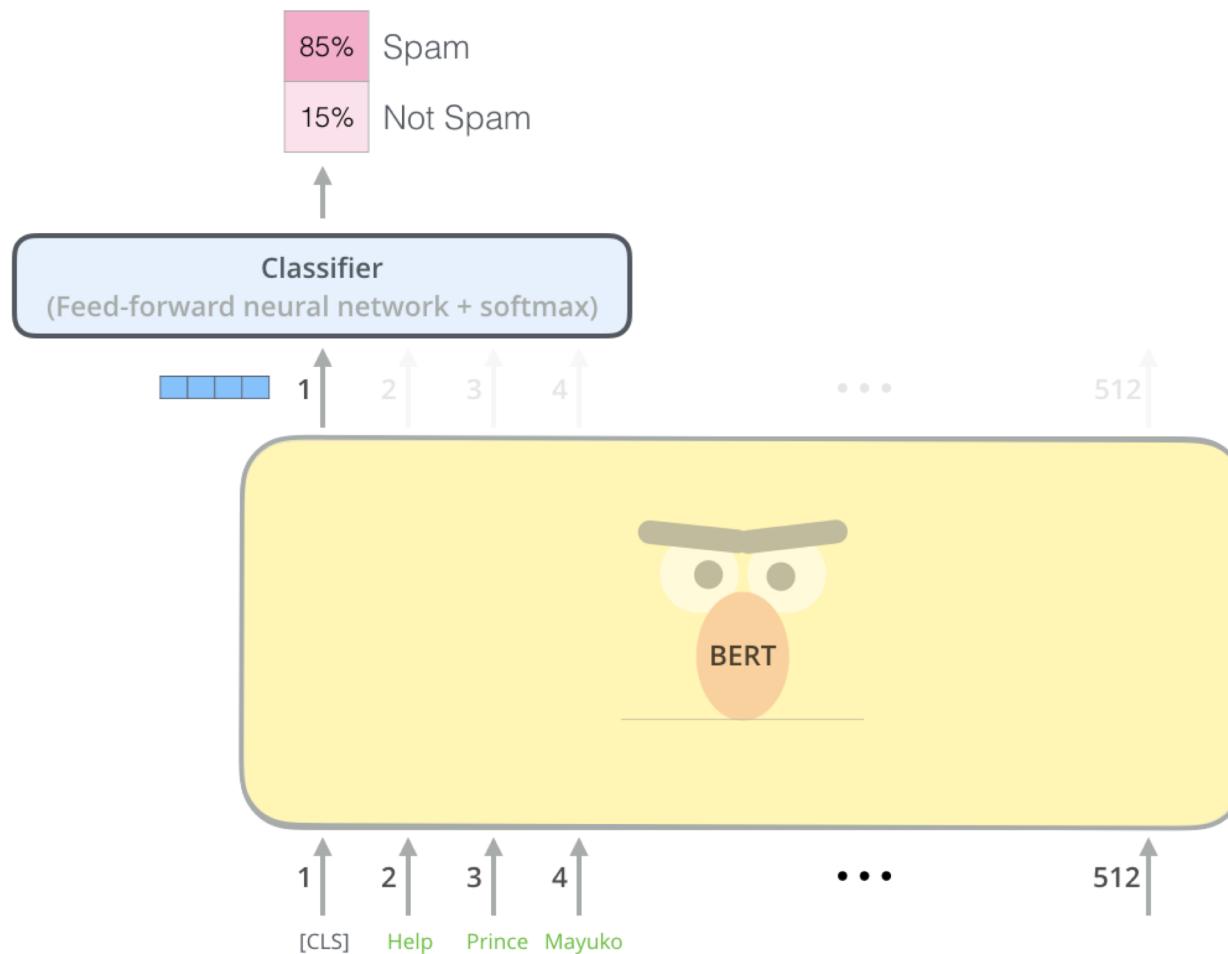
- A **model** to extract *contextualized* word embedding





# BERT: Pre-trained Text Representation Model

- Use BERT for sentence classification





# BERT Results

- Huge improvements over SOTA on 12 NLP task

System	MNLI-(m/mm)	QQP	QNLI	SST-2	CoLA	STS-B	MRPC	RTE	Average
	392k	363k	108k	67k	8.5k	5.7k	3.5k	2.5k	-
Pre-OpenAI SOTA	80.6/80.1	66.1	82.3	93.2	35.0	81.0	86.0	61.7	74.0
BiLSTM+ELMo+Attn	76.4/76.1	64.8	79.9	90.4	36.0	73.3	84.9	56.8	71.0
OpenAI GPT	82.1/81.4	70.3	88.1	91.3	45.4	80.0	82.3	56.0	75.2
BERT <sub>BASE</sub>	84.6/83.4	71.2	90.1	93.5	52.1	85.8	88.9	66.4	79.6
BERT <sub>LARGE</sub>	<b>86.7/85.9</b>	<b>72.1</b>	<b>91.1</b>	<b>94.9</b>	<b>60.5</b>	<b>86.5</b>	<b>89.3</b>	<b>70.1</b>	<b>81.9</b>

Table 1: GLUE Test results, scored by the GLUE evaluation server. The number below each task denotes the number of training examples. The “Average” column is slightly different than the official GLUE score, since we exclude the problematic WNLI set. OpenAI GPT = (L=12, H=768, A=12); BERT<sub>BASE</sub> = (L=12, H=768, A=12); BERT<sub>LARGE</sub> = (L=24, H=1024, A=16). BERT and OpenAI GPT are single-model, single task. All results obtained from <https://gluebenchmark.com/leaderboard> and <https://blog.openai.com/language-unsupervised/>.





# BERT: Pre-training Procedure

- Model architecture:
  - A big Transformer Encoder (240M free parameters)
- Dataset:
  - Wikipedia (2.5B words) + a collection of free ebooks (800M words)
- Training procedure
  - **masked language model** (masked LM)
    - Masks some percent of words from the input and has to reconstruct those words from context





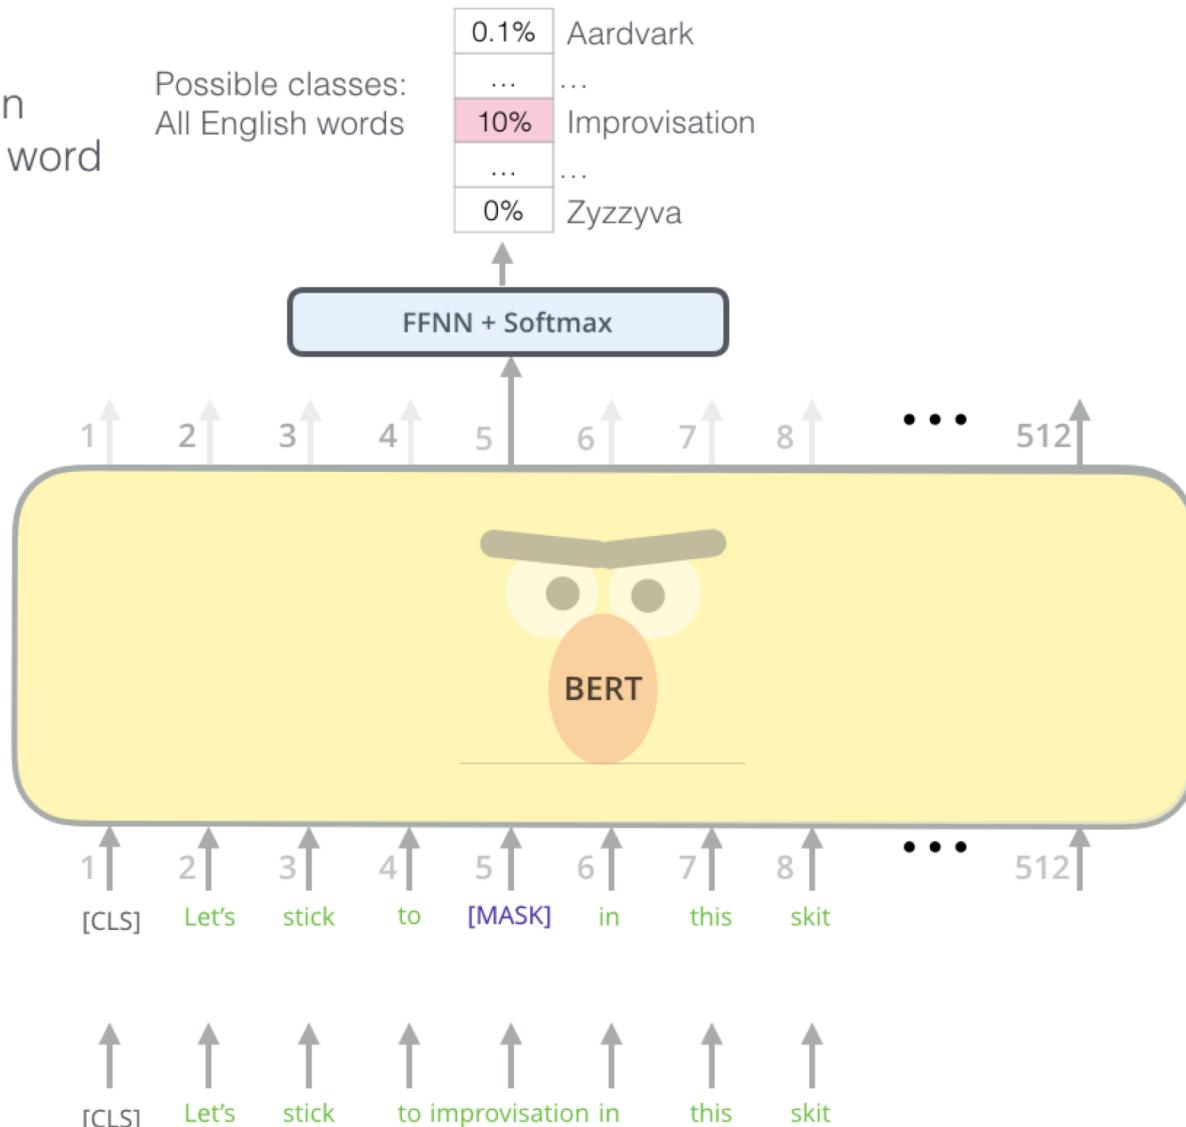
# BERT: Pre-training Procedure

- Masked LM

Use the output of the masked word's position to predict the masked word

Randomly mask 15% of tokens

Input





# BERT: Pre-training Procedure

- Model architecture:
  - A big Transformer Encoder (240M free parameters)
- Dataset:
  - Wikipedia (2.5B words) + a collection of free ebooks (800M words)
- Training procedure
  - **masked language model** (masked LM)
    - Masks some percent of words from the input and has to reconstruct those words from context
  - **Two-sentence task**
    - To understand relationships between sentences
    - Concatenate two sentences A and B and predict whether B actually comes after A in the original text





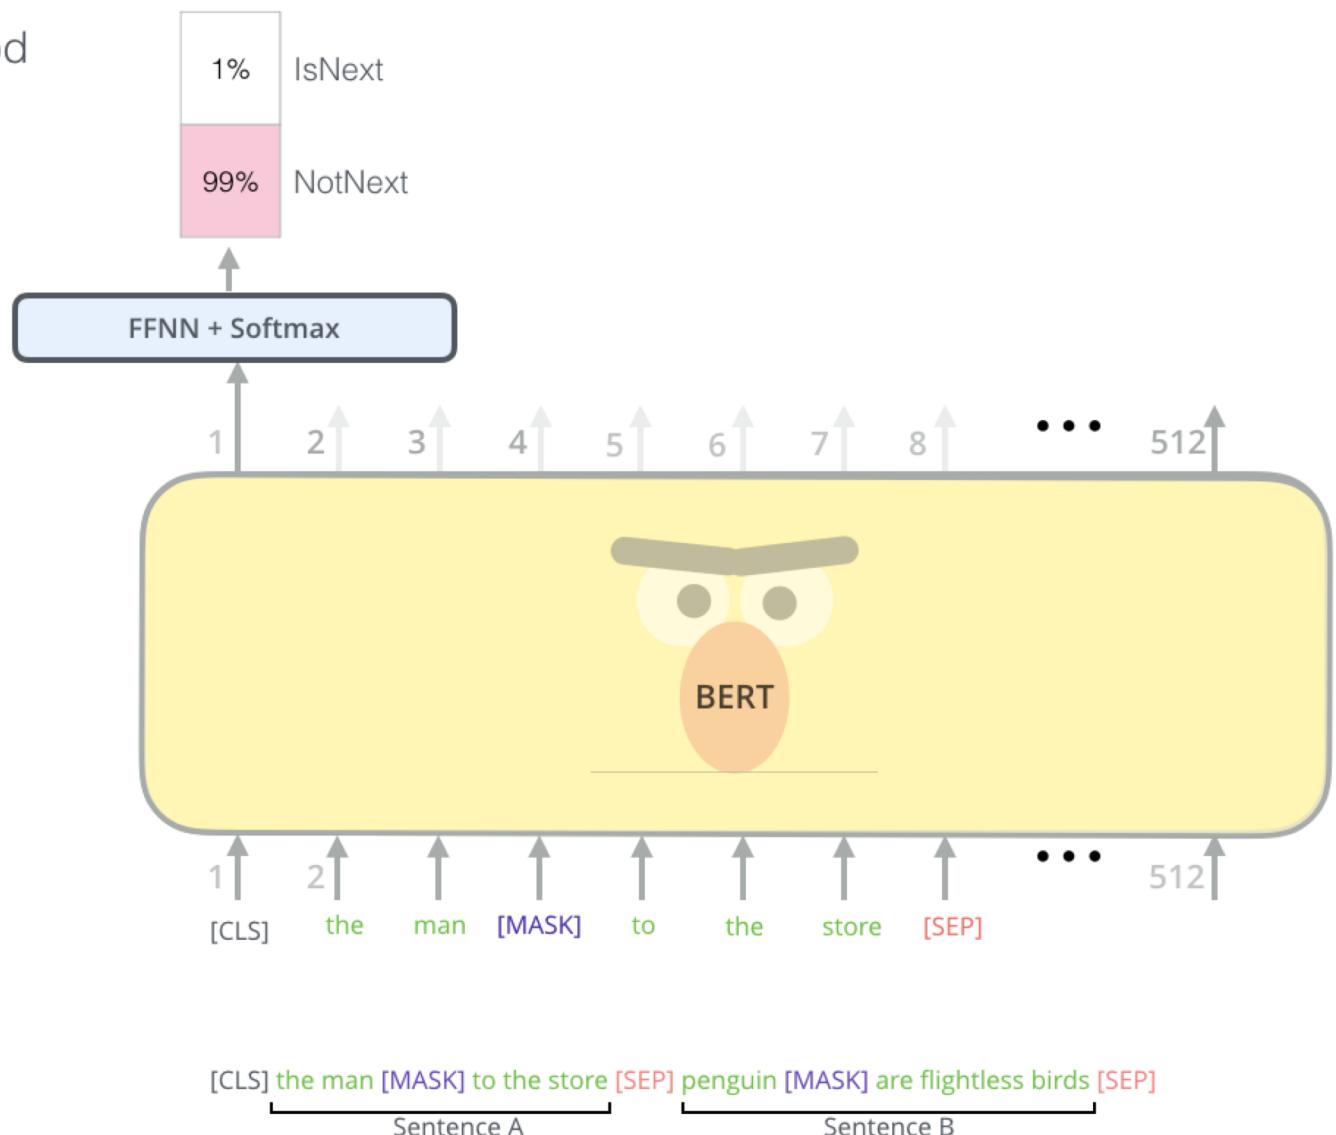
# BERT: Pre-training Procedure

- Two sentence task

Predict likelihood  
that sentence B  
belongs after  
sentence A

Tokenized  
Input

Input





# BERT: Pre-training Procedure

- BERT is trained on 4 TPU pods (=256 TPU chips) in 4 days
  - TPU: a matrix multiplication engine
- = 64 V100 GPUs, Infiniband network, 5.3 days
- = a standard 4 GPU desktop with RTX 2080Ti, 99 days



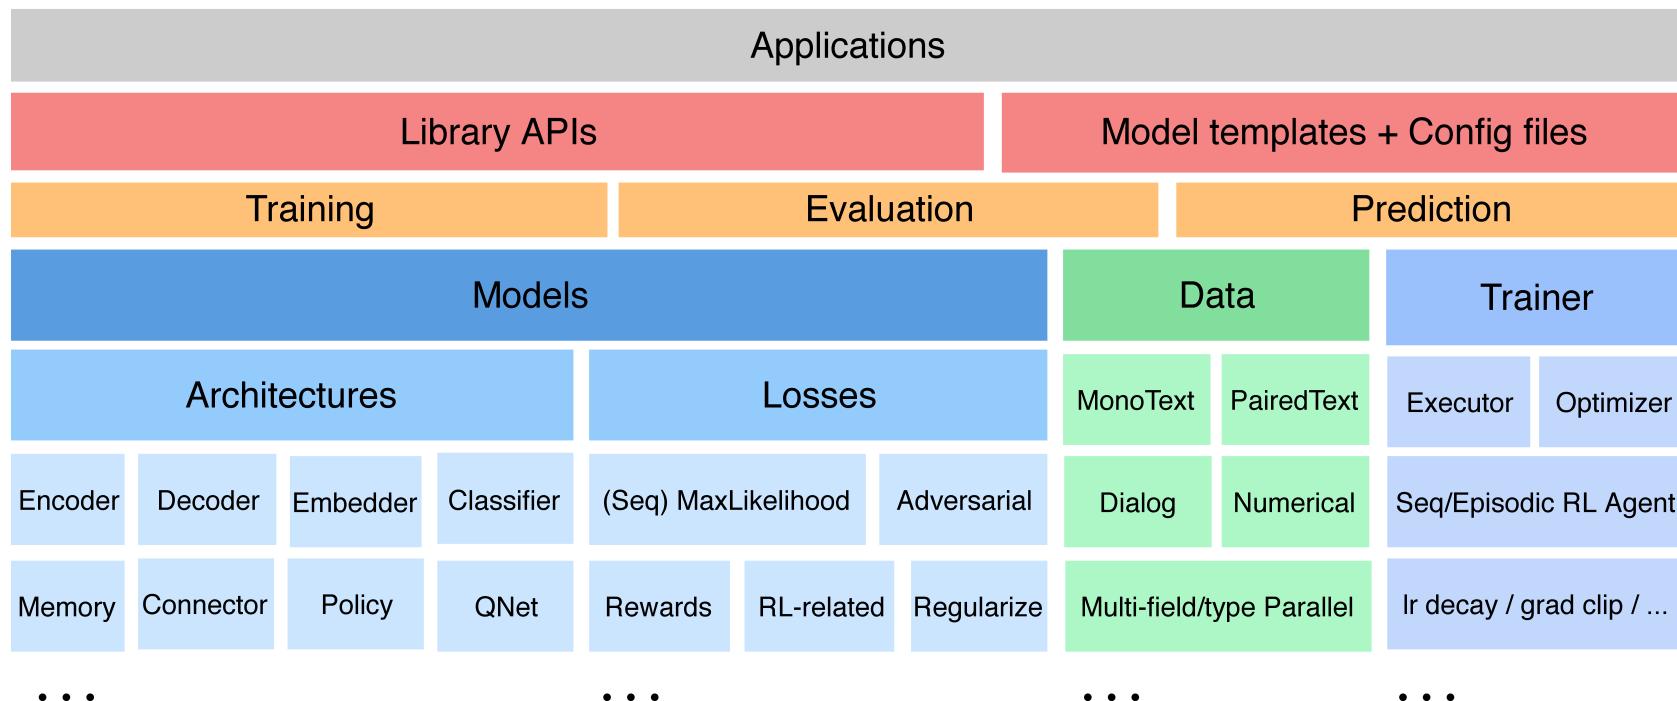


# Word Embedding on Texar



- A general-purpose text generation toolkit on TensorFlow

*Texar stack*





# Word Embedding on Texar

- Word2vec, Glove

```
1 import texar as tx
2
3 # Load data and pre-trained word embedding matrix
4 data = tx.data.MonoTextData(hparams=config.data)
5 iterator = tx.data.DataIterator(data)
6 data_batch = iterator.get_next()
7
8 # Create and initialize word embedder
9 embedder = texar.modules.WordEmbedder(
10     init_value=data.embedding_init_value, hparams=config.emb)
11
12 # Embed text into vectors
13 data_embed = embedder(data_batch)
14
15 # Downstream tasks
16 classifier = tx.modules.Conv1DClassifier(hparams=config.clas)
17 logits, pred = classifier(input=data_embed)
```

```
21 config.data = {
22     "embedding_init": {
23         "file": "word2vec.pretrain.dat"
24         "read_fn": "load_word2vec" # "load_glove"
25     }
26 }
```





# Word Embedding on Texar

- Word2vec, Glove

```
1 import texar as tx
2
3 # Load data and pre-trained word embedding matrix
4 data = tx.data.MonoTextData(hparams=config.data)
5 iterator = tx.data.DataIterator(data)
6 data_batch = iterator.get_next()
7
8 # Create and initialize word embedder
9 embedder = texar.modules.WordEmbedder(
10     init_value=data.embedding_init_value, hparams=co
11
12 # Embed text into vectors
13 data_embed = embedder(data_batch)
14
15 # Downstream tasks
16 classifier = tx.modules.Conv1DClassifier(hparams=config.clas)
17 logits, pred = classifier(input=data_embed)
```

```
29 # Create BERT embedder
30 embedder = tx.modules.TransformerEncoder(hparams=bert_config)
31 # Initialize BERT embedder
32 texar.init_bert_checkpoint("./bert.ckpt")
33
34 # Embed text into vectors
35 data_embed = embedder(data_batch)
```





# Seq2seq Attention on Texar

---

```
1 # Read data
2 dataset = PairedTextData(data_hparams)
3 batch = DataIterator(dataset).get_next()
4
5 # Encode
6 embedder = WordEmbedder(dataset.vocab.size, hparams=embedder_hparams)
7 encoder = TransformerEncoder(hparams=encoder_hparams)
8 enc_outputs = encoder(embedder(batch['source_text_ids']),
9                      batch['source_length'])
10
11 # Decode
12 decoder = AttentionRNNDecoder(memory=enc_outputs,
13                                 hparams=decoder_hparams)
14 outputs, length, _ = decoder(inputs=embedder(batch['target_text_ids']),
15                             seq_length=batch['target_length']-1)
16
17 # Loss
18 loss = sequence_sparse_softmax_cross_entropy(
19     labels=batch['target_text_ids'][:,1:], logits=outputs.logits, seq_length=length)
```

---





# Seq2seq Attention on Texar

```
1 # Read data
2 dataset = PairedTextData(data_hparams)
3 batch = Dataliterator(dataset).get_next()
4
5 # Encode
6 embedder = WordEmbedder(dataset.vocab.size, hparams=embedder_hparams)
7 encoder = TransformerEncoder(hparams=encoder_hparams)
8 enc_outputs = encoder(embedder(batch['source_text_ids']),
9                      batch['source_length'])
10
11 # Decode
12 decoder = AttentionRNNDecoder(memory=enc_outputs,
13                                 hparams=decoder_hparams)
14 outputs, length, _ = decoder(inputs=embedder(batch['target_text_ids']),
15                             seq_length=batch['target_length']-1)
16
17 # Loss
18 loss = sequence_sparse_softmax_cross_entropy(
19     labels=batch['target_text_ids'][:,1:], logits=outputs.logits, seq_length=length)
20
```

```
1 decoder_hparams = {
2     'rnn_cell': {
3         'type': 'LSTMCell'
4     }
5     'num_layers': 2,
6     'attention': {
7         'type': 'LuongAttention',
8         'kwargs': {
9             'num_units': 256,
10        }
11    }
12 }
```





# Takeaways

- Convolutional Networks (ConvNets)
- Recurrent Networks (RNNs)
  - LSTM designed for long-range dependency, vanishing
  - RNNs not only for sequence data, but also 2D sequences, Trees, graphs
- Attention Mechanisms
  - Three core elements: (Query, Key, Value)
  - Many variants based on alignment score functions
  - Attention on Text and Images
- Transformers: Multi-head Attention
  - Transformer: encoder-decoder
  - BERT: pre-trained text representation
  - GPT-2: pre-trained language model

